


Optimizing Infrastructure: Transitioning to Bare Metal Kubernetes

Jing Yan
Jul 13, 2024

Agenda

- Background
 - Evolution of infrastructure
 - Motivation towards bare metal k8s
- Bare metal k8s
 - Benefits
 - Challenges
 - Key Technologies
 - Metal³ and How Metal³ manages (bare metal) machines
 - **Cluster API (CAPI) and How CAPI works**
 - Demo time: CAPI
- Summary
- References



Background: evolution of infrastructure

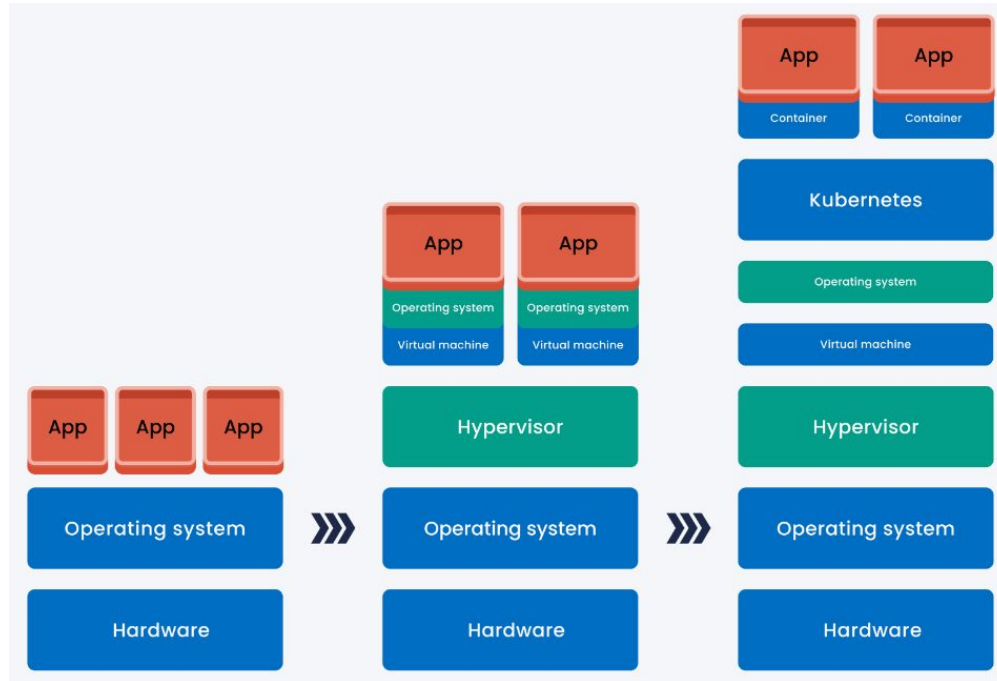


Figure: Evolution of Infrastructure [1]

Background: new needs sensitive to hypervisor overhead

Machine Learning as a Service (MLaaS)

- Training and Inference Tasks
- Real-Time Analytics

High-Performance Computing (HPC)

- Scientific Simulations
- Financial Modeling

Big Data Analytics

- Data Processing Pipelines
- Streaming Analytics

Real-Time Applications


- Online Gaming
- Augmented Reality (AR) and Virtual Reality (VR)

Database as a Service (DBaaS)

- High-Transaction Databases
- In-Memory Databases

Edge Computing

- IoT Devices
- Autonomous Vehicles



Bare metal k8s

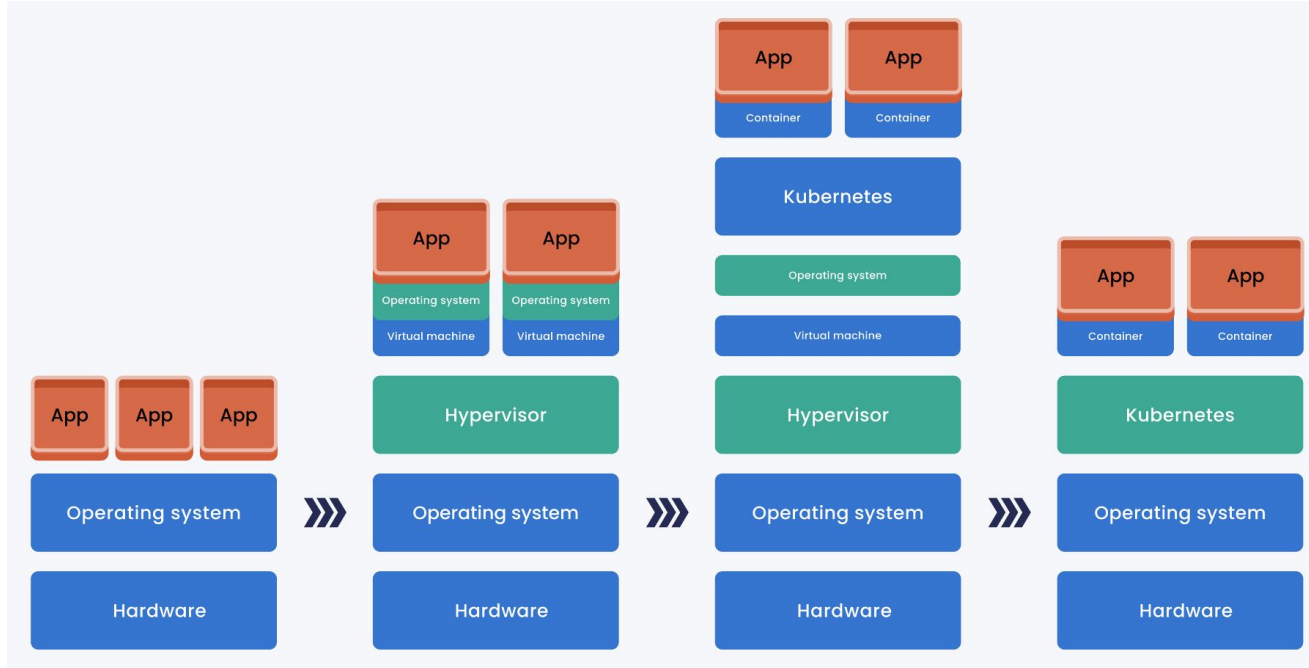


Figure: Evolution of Infrastructure [1]

Bare metal k8s - benefits

Deploying and Running k8s on top of Bare Metal Machines instead of Virtual Machines means one layer less to manage: the hypervisor.

Eliminating the hypervisor with bare metal servers can offer:

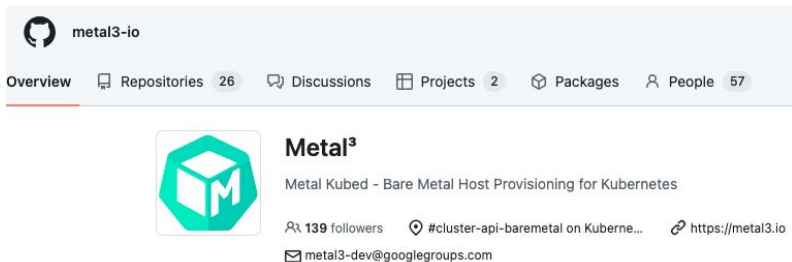
- Increased performance, particularly for specialized workloads like big data and AI.
- Better resource utilization, by as much as 20% [1].
- Savings, both in software licensing and operational efficiencies.
- Reduced risk through fewer technical and commercial dependencies.
- Better visibility into the health of your hardware.
- ...

Bare metal k8s - challenges

- How to provision a large number of bare metal machines
 - Solution: **Metal³** - bare metal server lifecycle management
- How to provision, upgrade, and operate multiple k8s clusters
 - Solution: **Cluster API (CAPI)** - cluster lifecycle management
- Configuration (Networking) Complexity
 - Compare with the relative simplicity of managing VM images, configuring **networking, storage**, and other resources in a bare metal Kubernetes environment provides more flexibility but can be more complex.

Metal³ [2][3]


Metal³ (pronounced “metal cubed”) is an open-source project that provides a set of tools for managing bare-metal infrastructure using Kubernetes [2].




The screenshot shows the GitHub repository page for metal3-io. At the top, the repository name 'metal3-io' is displayed with the GitHub logo. Below this, navigation tabs include 'Overview' (selected), 'Repositories' (26), 'Discussions', 'Projects' (2), 'Packages', and 'People' (57). The main content area features the 'Metal³' repository card. The card has a green cube icon with a white 'M' and the title 'Metal³'. Below the title is the description 'Metal Kubed - Bare Metal Host Provisioning for Kubernetes'. Further down, it shows '139 followers', a location pin icon for '#cluster-api-baremetal on Kuberne...', a link icon for 'https://metal3.io', and an email icon for 'metal3-dev@googlegroups.com'.



Ironic

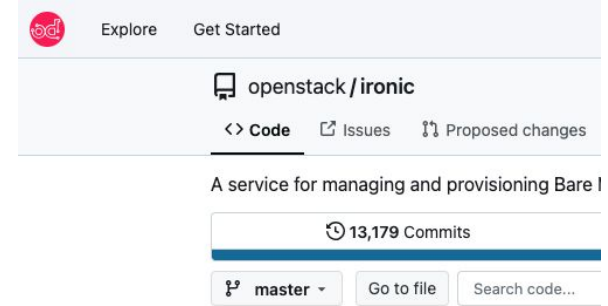
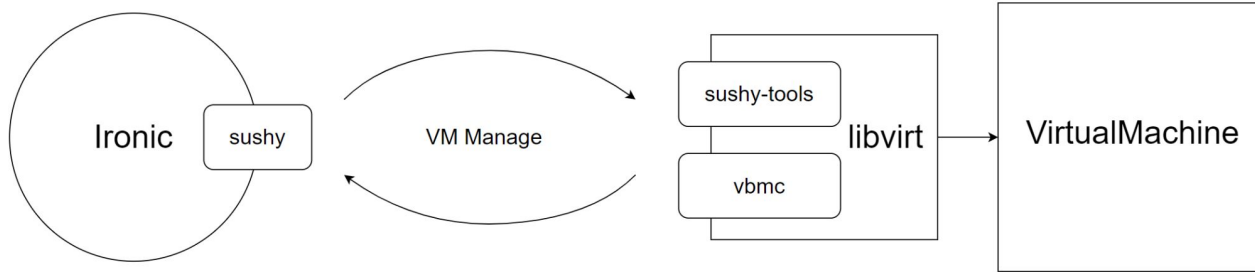
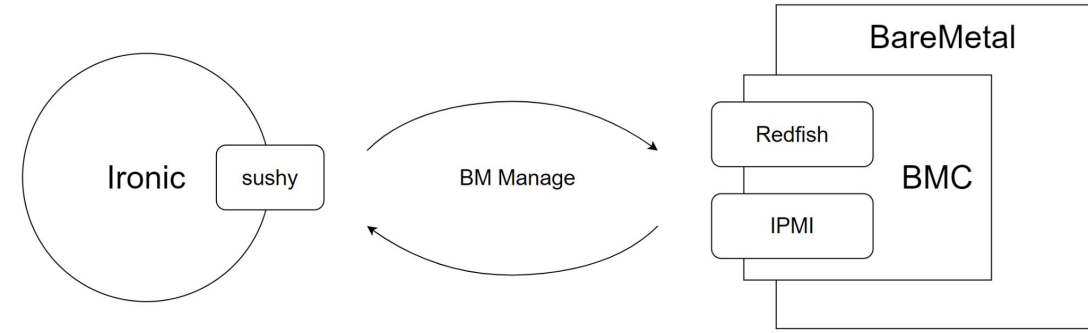


Bare Metal Operator



Cluster API Provider
Metal³

How Metal³ manages (bare metal) machines



IRONIC

an OpenStack Community Project

CAPI [4]




CAPI is a tool for programmatically configuring and deploying k8s clusters on a variety of different infrastructures.

- Declarative k8s management
 - CAPI enables you to describe and manage the lifecycle of clusters using YAML definitions for resources such as Cluster and Machine (**Like defining Pods and Services in Kubernetes using YAML files**).
- **Platform Diversity**
 - CAPI supports multiple infrastructure providers, including bare metal servers, public clouds, and private clouds.



CAPI: platform diversity



Provider Implementations

The code in this repository is independent of any specific deployment environment. Provider specific code is being developed in separate repositories, some of which are also sponsored by SIG Cluster Lifecycle. Check provider's documentation for updated info about which API version they are supporting.

Bootstrap

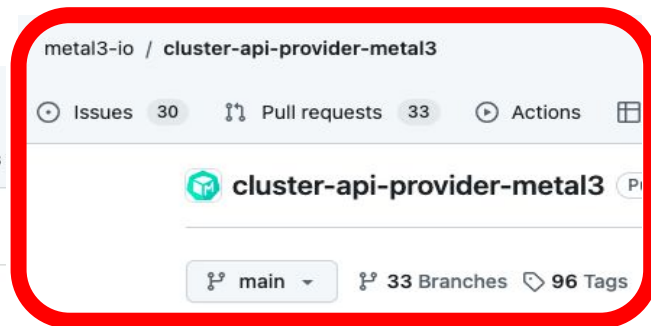




- Amazon Elastic Kubernetes Service (EKS)
- Kubeadm
- MicroK8s
- Oracle Cloud Native Environment (OCNE)
- Talos
- K3s
- k0smotron/k0s

Control Plane


- Kubeadm
- MicroK8s
- Nested
- Oracle Cloud Native Environment (OCNE)
- Talos
- Kamaji
- K3s
- k0smotron/k0s

Infrastructure

- Akamai (Linode)
- AWS



CAPI: deployment (1)



Deploying CAPI involves two k8s clusters: one is a temporary cluster called the **bootstrap cluster** that will be discarded later, which creates a second cluster that becomes the permanent CAPI **management cluster**.

Management cluster: responsible for creating and overseeing other clusters through Cluster API. They are your agents for infrastructure management, focusing entirely on provisioning, monitoring, and managing other clusters.

Workload cluster(s): handle application workloads for users, running microservices, and handling requests.

CAPI: deployment (2)

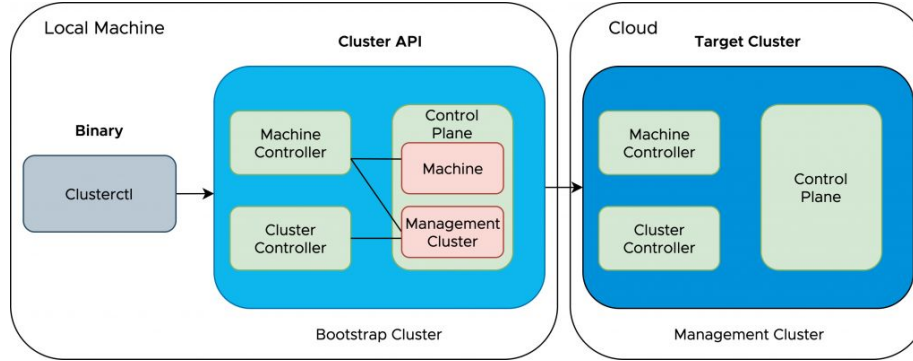


Figure: Deployment of Management Cluster [5]

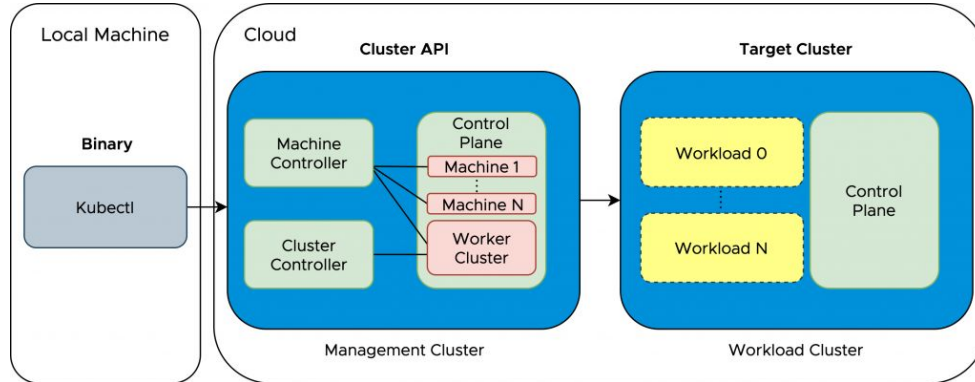


Figure: Deployment of Workload Cluster [5]

- Light blue: a Kubernetes in Docker (KinD) cluster.
- Dark blue: a cluster deployed to an infrastructure provider.
- Green: a controller or set of controllers.
- Red: CustomResourceDefinitions (CRDs).
- Yellow: user workloads.

CAPi: set up temporary bootstrap k8s cluster

```
jingyan@JingdeMBP clusterctl % kind create cluster --config kind-cluster-with-extramounts.yaml --name kind
Creating cluster "kind" ...
✓ Ensuring node image (kindest/node:v1.30.0)
✓ Preparing nodes
✓ Writing configuration
✓ Starting control-plane
✓ Installing CNI
✓ Installing StorageClass
Set kubectl context to "kind-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-kind

Thanks for using kind!

jingyan@JingdeMBP clusterctl % docker container ps -a
CONTAINER ID   IMAGE                                COMMAND                  CREATED        STATUS        PORTS                               NAMES
19c9fa083c22   kindest/node:v1.30.0               "/usr/local/bin/entr..." About a minute ago Up About a minute 127.0.0.1:49935->6443/tcp   kind-control-plane

jingyan@JingdeMBP clusterctl %
jingyan@JingdeMBP clusterctl % kubectl get pods -A
NAMESPACE      NAME                                      READY   STATUS    RESTARTS   AGE
kube-system    coredns-7db6d8ff4d-2mnpb               1/1     Running   0          63s
kube-system    coredns-7db6d8ff4d-p4vsv               1/1     Running   0          63s
kube-system    etcd-kind-control-plane                 1/1     Running   0          78s
kube-system    kindnet-hhvg4                           1/1     Running   0          63s
kube-system    kube-apiserver-kind-control-plane       1/1     Running   0          78s
kube-system    kube-controller-manager-kind-control-plane 1/1     Running   0          78s
kube-system    kube-proxy-xc9hn                       1/1     Running   0          63s
kube-system    kube-scheduler-kind-control-plane       1/1     Running   0          78s
local-path-storage local-path-provisioner-988d74bc-k7js9  1/1     Running   0          63s

jingyan@JingdeMBP clusterctl %
jingyan@JingdeMBP clusterctl % kubectl config current-context
kind-kind
```

CAPI: set up management k8s cluster

```
jingyan@JingdeMBP clusterctl % clusterctl init --infrastructure docker
Fetching providers
Installing cert-manager Version="v1.15.1"
Waiting for cert-manager to be available...
Installing Provider="cluster-api" Version="v1.7.4" TargetNamespace="capi-system"
Installing Provider="bootstrap-kubeadm" Version="v1.7.4" TargetNamespace="capi-kubeadm-bootstrap-system"
Installing Provider="control-plane-kubeadm" Version="v1.7.4" TargetNamespace="capi-kubeadm-control-plane-system"
Installing Provider="infrastructure-docker" Version="v1.7.4" TargetNamespace="capd-system"

Your management cluster has been initialized successfully!
```

```
jingyan@JingdeMacBook-Pro clusterctl % kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
capd-system	capd-controller-manager-bfb455d6d-tpkdz	1/1	Running	0	10m
capi-kubeadm-bootstrap-system	capi-kubeadm-bootstrap-controller-manager-6868fcb86f-td7dj	1/1	Running	0	10m
capi-kubeadm-control-plane-system	capi-kubeadm-control-plane-controller-manager-7466b4f659-jklft	1/1	Running	0	10m
capi-system	capi-controller-manager-79949bb88f-pt52h	1/1	Running	0	10m
cert-manager	cert-manager-cainjector-9d956987c-5f8tf	1/1	Running	0	11m
cert-manager	cert-manager-fdd97855b-hdsf5	1/1	Running	0	11m
cert-manager	cert-manager-webhook-9f799c7d7-999gs	1/1	Running	0	11m
kube-system	coredns-7db6d8ff4d-2mnpb	1/1	Running	0	22m
kube-system	coredns-7db6d8ff4d-p4vsv	1/1	Running	0	22m
kube-system	etcd-kind-control-plane	1/1	Running	0	22m
kube-system	kindnet-hhvg4	1/1	Running	0	22m
kube-system	kube-apiserver-kind-control-plane	1/1	Running	0	22m
kube-system	kube-controller-manager-kind-control-plane	1/1	Running	0	22m
kube-system	kube-proxy-xc9hn	1/1	Running	0	22m
kube-system	kube-scheduler-kind-control-plane	1/1	Running	0	22m
local-path-storage	local-path-provisioner-988d74bc-k7js9	1/1	Running	0	22m

CAPi: set up workload k8s clusters (1)

```
jingyan@JingdeMBP clusterctl % clusterctl generate cluster workload0 --from ./cluster-template-development.yaml \
--kubernetes-version v1.30.0 \
--control-plane-machine-count=3 \
--worker-machine-count=3 \
> workload0.yaml
jingyan@JingdeMBP clusterctl %
jingyan@JingdeMBP clusterctl % cat workload0.yaml
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
metadata:
  name: workload0
  namespace: default
spec:
  clusterNetwork:
    pods:
      cidrBlocks:
        - 192.168.0.0/16
      serviceDomain: cluster.local
    services:
      cidrBlocks:
        - 10.128.0.0/12
  topology:
    class: quick-start
    controlPlane:
      metadata: {}
      replicas: 3
    variables:
      - name: imageRepository
        value: ""
      - name: etcdImageTag
        value: ""
      - name: coreDNSImageTag
        value: ""
      - name: podSecurityStandard
        value:
          audit: restricted
          enabled: false
          enforce: baseline
          warn: restricted
    version: v1.30.0
  workers:
    machineDeployments:
      - class: default-worker
        name: md-0
        replicas: 3
    machinePools:
      - class: default-worker
        name: mp-0
        replicas: 3
```

```
jingyan@JingdeMBP clusterctl % kubectl apply -f workload0.yaml
cluster.cluster.x-k8s.io/workload0 configured
jingyan@JingdeMBP clusterctl %
jingyan@JingdeMBP clusterctl % kubectl get clusters -A -o wide
```

NAMESPACE	NAME	CLUSTERCLASS	PHASE	AGE	VERSION
default	workload0	quick-start	Provisioned	3m26s	v1.30.0

```
jingyan@JingdeMBP clusterctl %
```

```
jingyan@JingdeMBP clusterctl % kubectl get machines -A -o wide
```

NAMESPACE	NAME	CLUSTER	NODENAME	PROVIDERID	PHASE	AGE	VERSION
default	worker-b1lhd3	workload0	workload0-worker-b1lhd3	dacker:///workload0-worker-b1lhd3	Running	84s	
default	worker-l8o58j	workload0	workload0-worker-l8o58j	dacker:///workload0-worker-l8o58j	Running	83s	
default	worker-rza4os	workload0	workload0-worker-rza4os	dacker:///workload0-worker-rza4os	Running	83s	
default	workload0-4gc9s-klmd8	workload0	workload0-4gc9s-klmd8	dacker:///workload0-4gc9s-klmd8	Running	106s	v1.30.0
default	workload0-4gc9s-nhwmz	workload0	workload0-4gc9s-nhwmz	dacker:///workload0-4gc9s-nhwmz	Running	58s	v1.30.0
default	workload0-4gc9s-pnpxw	workload0	workload0-4gc9s-pnpxw	dacker:///workload0-4gc9s-pnpxw	Running	2m37s	v1.30.0
default	workload0-md-0-44x8t-j2c2t-5mvlv	workload0	workload0-md-0-44x8t-j2c2t-5mvlv	dacker:///workload0-md-0-44x8t-j2c2t-5mvlv	Running	2m56s	v1.30.0
default	workload0-md-0-44x8t-j2c2t-njpc2	workload0	workload0-md-0-44x8t-j2c2t-njpc2	dacker:///workload0-md-0-44x8t-j2c2t-njpc2	Running	2m56s	v1.30.0
default	workload0-md-0-44x8t-j2c2t-wvlh7	workload0	workload0-md-0-44x8t-j2c2t-wvlh7	dacker:///workload0-md-0-44x8t-j2c2t-wvlh7	Running	2m56s	v1.30.0

CAPI: set up workload k8s clusters (2)

```
jingyan@JingdeMBP clusterctl % clusterctl generate cluster workload1 --from ./cluster-template-development.yaml \
--kubernetes-version v1.27.0 \
--control-plane-machine-count=1 \
--worker-machine-count=1 \
> workload1.yaml

jingyan@JingdeMBP clusterctl % kubectl apply -f workload1.yaml
cluster.cluster.x-k8s.io/workload1 configured

jingyan@JingdeMBP clusterctl %
jingyan@JingdeMBP clusterctl % kubectl get clusters -A


| NAMESPACE | NAME      | CLUSTERCLASS | PHASE       | AGE | VERSION |
|-----------|-----------|--------------|-------------|-----|---------|
| default   | workload0 | quick-start  | Provisioned | 32m | v1.30.0 |
| workload1 | workload1 | quick-start  | Provisioned | 83s | v1.27.0 |



jingyan@JingdeMBP clusterctl %
jingyan@JingdeMBP clusterctl % kubectl get machines -n workload1


| NAME                             | CLUSTER   | NODENAME                         | PROVIDERID                                 | PHASE   | AGE | VERSION |
|----------------------------------|-----------|----------------------------------|--------------------------------------------|---------|-----|---------|
| worker-c9xbnw                    | workload1 | workload1-worker-c9xbnw          | docker:///workload1-worker-c9xbnw          | Running | 24s |         |
| workload1-dps2s-7qzj6            | workload1 | workload1-dps2s-7qzj6            | docker:///workload1-dps2s-7qzj6            | Running | 88s | v1.27.0 |
| workload1-md-0-plqh5-r2qd9-7zhjl | workload1 | workload1-md-0-plqh5-r2qd9-7zhjl | docker:///workload1-md-0-plqh5-r2qd9-7zhjl | Running | 90s | v1.27.0 |


```

```
jingyan@JingdeMBP clusterctl % kubectl describe cluster workload0 | grep "Cluster Network" -A10
```

```
Cluster Network:
  Pods:
    Cidr Blocks:
      192.168.0.0/16
    Service Domain: cluster.local
    Services:
      Cidr Blocks:
        10.128.0.0/12
    Control Plane Endpoint:
      Host: 172.18.0.3
      Port: 6443
```

```
jingyan@JingdeMBP clusterctl % kubectl describe cluster workload1 -n workload1 | grep "Cluster Network" -A10
Cluster Network:
  Pods:
    Cidr Blocks:
      192.168.0.0/16
    Service Domain: cluster.local
    Services:
      Cidr Blocks:
        10.128.0.0/12
    Control Plane Endpoint:
      Host: 172.18.0.13
      Port: 6443
```

CAPI: deploy cni for workload0 (1)

```
jingyan@JingdeMacBook-Pro clusterctl % clusterctl get kubeconfig workload0 > workload0.kubeconfig
```

```
root@kind-control-plane:~# kubectl --kubeconfig=./workload0.kubeconfig get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
workload0-4gc9s-klmd8	NotReady	control-plane	42m	v1.30.0
workload0-4gc9s-nhwwz	NotReady	control-plane	42m	v1.30.0
workload0-4gc9s-pnpxw	NotReady	control-plane	43m	v1.30.0
workload0-md-0-44x8t-j2c2t-5mvlv	NotReady	<none>	42m	v1.30.0
workload0-md-0-44x8t-j2c2t-njpc2	NotReady	<none>	42m	v1.30.0
workload0-md-0-44x8t-j2c2t-wwlh7	NotReady	<none>	42m	v1.30.0
workload0-worker-b1lhd3	NotReady	<none>		
workload0-worker-l8o58j	NotReady	<none>		
workload0-worker-rza4os	NotReady	<none>		

```
root@kind-control-plane:~# kubectl --kubeconfig=./workload0.kubeconfig apply -f https://raw.githubusercontent.com/projectcalico/calico/v3.26.1/manifests/calico.yaml
poddissruptionbudget.policy/calico-kube-controllers created
serviceaccount/calico-kube-controllers created
serviceaccount/calico-node created
serviceaccount/calico-cni-plugin created
configmap/calico-config created
customresourcedefinition.apiextensions.k8s.io/bgppconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgppfilters.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/bgpppeers.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/blockaffinities.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/caliconodestatuses.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/clusterinformations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/felixconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/globalnetworksets.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/hostendpoints.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamblocks.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamconfigs.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/ipamhandles.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/lppools.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/lpreservations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/kubecontrollersconfigurations.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networkpolicies.crd.projectcalico.org created
customresourcedefinition.apiextensions.k8s.io/networksets.crd.projectcalico.org created
clusterrole.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrole.rbac.authorization.k8s.io/calico-node created
clusterrole.rbac.authorization.k8s.io/calico-cni-plugin created
clusterrolebinding.rbac.authorization.k8s.io/calico-kube-controllers created
clusterrolebinding.rbac.authorization.k8s.io/calico-node created
clusterrolebinding.rbac.authorization.k8s.io/calico-cni-plugin created
daemonset.apps/calico-node created
deployment.apps/calico-kube-controllers created
```

CAPI: deploy cni for workload0 (2)

```
root@kind-control-plane:/# kubectl --kubeconfig=./workload0.kubeconfig get nodes -A
```

NAME	STATUS	ROLES	AGE	VERSION
workload0-4gc9s-klmd8	Ready	control-plane	87m	v1.30.0
workload0-4gc9s-nhwwz	Ready	control-plane	87m	v1.30.0
workload0-4gc9s-pnpwx	Ready	control-plane	88m	v1.30.0
workload0-md-0-44x8t-j2c2t-5mvlv	Ready	<none>	87m	v1.30.0
workload0-md-0-44x8t-j2c2t-njpc2	Ready	<none>	87m	v1.30.0
workload0-md-0-44x8t-j2c2t-wwlh7	Ready	<none>	87m	v1.30.0
workload0-worker-b1lh3	Ready	<none>	87m	v1.30.0
workload0-worker-l8o58j	Ready	<none>	87m	v1.30.0
workload0-worker-rza4os	Ready	<none>	87m	v1.30.0

```
root@kind-control-plane:/# kubectl --kubeconfig=./workload0.kubeconfig get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	calico-kube-controllers-7dc5458bc6-pxxrn	1/1	Running	0	42m
kube-system	calico-node-4zcl5	1/1	Running	0	42m
kube-system	calico-node-6xlvr	1/1	Running	0	42m
kube-system	calico-node-85q7t	1/1	Running	0	42m
kube-system	calico-node-89jnv	1/1	Running	0	42m
kube-system	calico-node-dclap	1/1	Running	0	42m
kube-system	calico-node-dldbp	1/1	Running	0	42m
kube-system	calico-node-fsgqp	1/1	Running	0	42m
kube-system	calico-node-hv86k	1/1	Running	0	42m
kube-system	calico-node-trshn	1/1	Running	0	42m
kube-system	coredns-7db6d8ff4d-g4p8b	1/1	Running	0	88m
kube-system	coredns-7db6d8ff4d-vnhxz	1/1	Running	0	88m
kube-system	etcd-workload0-4gc9s-klmd8	1/1	Running	0	87m
kube-system	etcd-workload0-4gc9s-nhwwz	1/1	Running	0	87m
kube-system	etcd-workload0-4gc9s-pnpwx	1/1	Running	0	88m
kube-system	kube-apiserver-workload0-4gc9s-klmd8	1/1	Running	0	87m
kube-system	kube-apiserver-workload0-4gc9s-nhwwz	1/1	Running	0	87m
kube-system	kube-apiserver-workload0-4gc9s-pnpwx	1/1	Running	0	88m
kube-system	kube-controller-manager-workload0-4gc9s-klmd8	1/1	Running	0	87m
kube-system	kube-controller-manager-workload0-4gc9s-nhwwz	1/1	Running	0	87m
kube-system	kube-controller-manager-workload0-4gc9s-pnpwx	1/1	Running	0	88m
kube-system	kube-proxy-g4cjf	1/1	Running	0	87m
kube-system	kube-proxy-h7v56	1/1	Running	0	87m
kube-system	kube-proxy-hnhxx	1/1	Running	0	87m
kube-system	kube-proxy-pwfcg	1/1	Running	0	87m
kube-system	kube-proxy-qfcdw	1/1	Running	0	87m
kube-system	kube-proxy-rlnfr	1/1	Running	0	87m
kube-system	kube-proxy-t62m2	1/1	Running	0	88m
kube-system	kube-proxy-tmfft	1/1	Running	0	87m
kube-system	kube-proxy-xskbd	1/1	Running	0	87m
kube-system	kube-scheduler-workload0-4gc9s-klmd8	1/1	Running	0	87m
kube-system	kube-scheduler-workload0-4gc9s-nhwwz	1/1	Running	0	87m
kube-system	kube-scheduler-workload0-4gc9s-pnpwx	1/1	Running	0	88m

CAPi: deploy cni for workload1

```
jingyan@JingdeMacBook-Pro clusterctl % clusterctl get kubeconfig workload1 -n workload1 > workload1.kubeconfig
```

```
root@kind-control-plane:/# kubectl --kubeconfig=./workload1.kubeconfig get nodes
```

NAME	STATUS	ROLES	AGE	VERSION
workload1-dps2s-7qzj6	NotReady	control-plane	13m	v1.27.0
workload1-md-0-plqh5-r2qd9-72hjl	NotReady	<none>	13m	v1.27.0
workload1-worker-c9xbnw	NotReady	<none>	13m	v1.27.0

```
root@kind-control-plane:/# kubectl --kubeconfig=./workload1.kubeconfig apply -f https://github.com/flannel-io/flannel/releases/latest/download/kube-flannel.yml
namespace/kube-flannel created
serviceaccount/flannel created
clusterrole.rbac.authorization.k8s.io/flannel created
clusterrolebinding.rbac.authorization.k8s.io/flannel created
configmap/kube-flannel-cfg created
daemonset.apps/kube-flannel-ds created
```

```
root@kind-control-plane:/# kubectl --kubeconfig=./workload1.kubeconfig get nodes -A
```

NAME	STATUS	ROLES	AGE	VERSION
workload1-dps2s-7qzj6	Ready	control-plane	58m	v1.27.0
workload1-md-0-plqh5-r2qd9-72hjl	Ready	<none>	58m	v1.27.0
workload1-worker-c9xbnw	Ready	<none>	58m	v1.27.0

CAPi: deploy service and access service

```
root@workload0-4gc9s-pnpwx:/# kubectl get deployments
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
nginx-deployment  3/3    3           3          10m

root@workload0-4gc9s-pnpwx:/# kubectl get services
NAME          TYPE        CLUSTER-IP    EXTERNAL-IP  PORT(S)
kubernetes    ClusterIP   10.128.0.1     <none>       443/TCP
nginx-service  NodePort    10.130.197.100 <none>       80:30000/
```

```
root@workload1-dps2s-7qzj6:/# wget -qO- curl http://172.18.0.4:30000
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

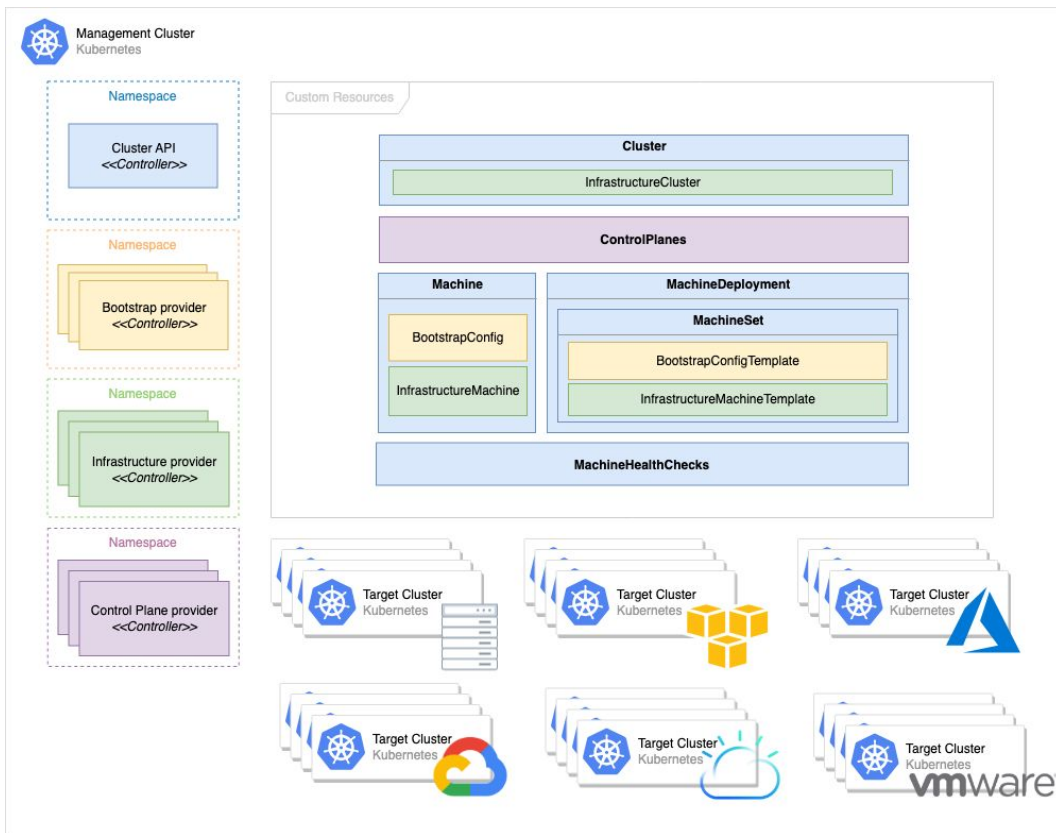
```
root@workload1-dps2s-7qzj6:/# kubectl get pods
NAME        READY  STATUS   RESTARTS  AGE
busybox     1/1    Running  0          2m9s

root@workload1-dps2s-7qzj6:/# kubectl exec -it busybox -- wget -qO- http://172.18.0.4:30000
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

How CAPI works [4]




Cluster API Controllers: responsible for managing the lifecycle of the Kubernetes clusters created using the Cluster API. These controllers are responsible for provisioning, scaling, and deleting the clusters, and ensuring that the Workload clusters are in the desired state.

Bootstrap provider: responsible for bootstrapping (installing and configuring) the Kubernetes control plane components on a newly created cluster.

Infrastructure provider: responsible for the provisioning of infrastructure resources required by the Cluster or by Machines.

Custom Resources: Kubernetes resources used by the Cluster API to create, manage, and delete clusters. These resources include **Cluster**, **Machine**, **MachineSet**, and **MachineDeployment**.

Summary



cluster-api: <https://github.com/kubernetes-sigs/cluster-api>

cluster-api-provider-metal3: <https://github.com/metal3-io/cluster-api-provider-metal3>

baremetal-operator: <https://github.com/metal3-io/baremetal-operator/tree/main>

ironic: <https://github.com/metal3-io/ironic-image>

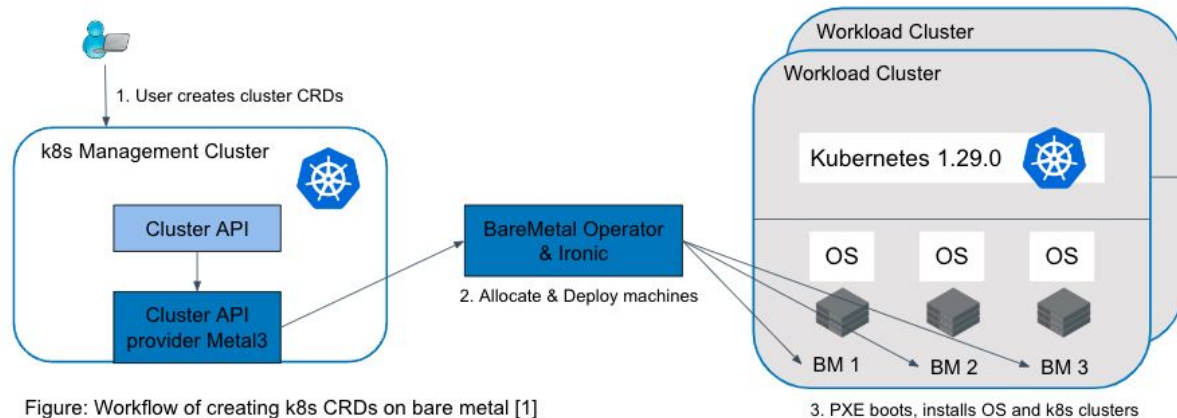


Figure: Workflow of creating k8s CRDs on bare metal [1]

References

- [1] <https://www.spectrocloud.com/blog/introducing-bare-metal-kubernetes-what-you-need-to-know>
- [2] <https://metal3.io/>
- [3] <https://github.com/metal3-io>
- [4] <https://cluster-api.sigs.k8s.io/introduction>
- [5] <https://tanzu.vmware.com/content/blog/pattern-recognition-how-cluster-api-reveals-the-core-of-kubernetes>