

One Cluster Multiple Providers

Date: Jul 27th, 2024

Author: Jing Yan

This document aims to provide a step-by-step hands-on tutorial for setting up Kubernetes clusters via [Cluster API](#) with different providers, namely [cluster-api-provider-metal3](#) and [cluster-api-provider-kubevirt](#), so that Kubernetes clusters could be set up on both bare metal machines and KubeVirt VMs as what I have shared before in the following two tech sharings.

Bare Metal Kubernetes: <https://drive.weixin.qq.com/s?k=ABAAngdnAAozJ5C4pz>

Metal3: <https://drive.weixin.qq.com/s?k=ABAAngdnAAoD1cTeyN>

And, there are some other components mentioned in this doc, feel free to browse for more detailed information.

Kind: <https://kind.sigs.k8s.io/>

Cluster API: <https://cluster-api.sigs.k8s.io/>; <https://github.com/kubernetes-sigs/cluster-api>

Calico: <https://docs.tigera.io/calico/latest/about/>; <https://github.com/projectcalico/calico>

Cert Manager: <https://cert-manager.io/>; <https://github.com/cert-manager/cert-manager>

Virtualization on Ubuntu: <https://ubuntu.com/server/docs/virtualisation-with-qemu>

Metal3: <https://book.metal3.io/introduction>; <https://github.com/metal3-io>

Bare Metal Operator: <https://book.metal3.io/bmo/introduction.html>;
<https://github.com/metal3-io/baremetal-operator>

MetalLB: <https://metallb.universe.tf/>; <https://github.com/metallb/metallb>

KubeVirt: <https://kubevirt-manager.io/>; <https://github.com/kubevirt>

cluster-api-provider-kubevirt: <https://github.com/kubernetes-sigs/cluster-api-provider-kubevirt>

cluster-api-provider-metal3: <https://github.com/metal3-io/cluster-api-provider-metal3>

Table of Contents

0. Environment description	3
1. Set up the bootstrap Kubernetes cluster via kind	4
2. Deploy Calico CNI	6
3. Deploy Cert Manager	6
4. Start VMs	7
5. Set up sushy-tools to manage VMs	8
6. Set up DHCP Server for VMs	9
7. Set up Image Server for VMs	9
8. Deploy Ironic for Metal³	10
9. Deploy Bare Metal Operator for Metal³	12
10. Create BareMetalHosts	13
11. Install MetalLB	14
12. Deploy KubeVirt	14
13. Initialize the management cluster	15
14. Create the target cluster on KubeVirt VMs only	16
15. Create the target cluster on Bare Metal Machines only (problematic)	17
16. Create the target cluster on Bare Metal Machines and KubeVirt VMs (problematic)	23
Conclusion	26

0. Environment description

In the following context, everything is deployed on this “X1 Carbon ThinkPad” laptop which has ubuntu-22.04 installed. If you would like to set the whole thing up on a distributed system, feel free to adjust the deployments.

```
# uname -a
Linux kubevirt-ThinkPad-X1-Carbon-Gen-11 6.5.0-1027-oem #28-Ubuntu SMP
PREEMPT_DYNAMIC Thu Jul 25 13:32:46 UTC 2024 x86_64 x86_64 x86_64 GNU/Linux
```

Docker will be necessary for the whole deployment, so just install it. Meanwhile, get your own “aliyun registry mirrors” and modify the docker configuration file accordingly IF YOU DO NOT HAVE A VPN && YOUR LOCATION IS CHINESE MAINLAND (which will save you ass thousands of times!!!).

```
# apt-get update
# apt-get install apt-transport-https ca-certificates curl
software-properties-common lrzsz -y
# curl -fsSL https://mirrors.aliyun.com/docker-ce/linux/ubuntu/gpg | apt-key add -
# add-apt-repository "deb [arch=amd64]
https://mirrors.aliyun.com/docker-ce/linux/ubuntu $(lsb_release -cs) stable"
# apt-get update
# apt-get install docker-ce docker-ce-cli containerd.io docker-buildx-plugin
docker-compose-plugin -y
# systemctl enable --now docker
# mkdir -p /etc/docker
# sudo tee /etc/docker/daemon.json <<-'EOF'
{
  "registry-mirrors": ["https://6pra95pl.mirror.aliyuncs.com"]
}
EOF
# systemctl daemon-reload
# systemctl restart docker
# systemctl status docker
```

You will need “kubectl” to access the Kubernetes clusters later, so install it as follows.

```
# apt-get update && apt-get install -y apt-transport-https
# curl -fsSL
https://mirrors.aliyun.com/kubernetes-new/core/stable/v1.30/deb/Release.key | gpg
--dearmor -o /etc/apt/keyrings/kubernetes-apt-keyring.gpg
# echo "deb [signed-by=/etc/apt/keyrings/kubernetes-apt-keyring.gpg]
https://mirrors.aliyun.com/kubernetes-new/core/stable/v1.30/deb/ /" | tee
/etc/apt/sources.list.d/kubernetes.list
# apt-get update
# apt-get install -y kubectl
# kubectl version
```

You will need qemu-kvm-based virtual machines (VMs) to simulate bare metal machines, so the virtualization components are also necessary, as well as enabling Virtualization in BIOS.

```
# apt install cpu-checker -y
# kvm-ok
INFO: /dev/kvm exists
KVM acceleration can be used
# apt -y install bridge-utils cpu-checker libvirt-clients libvirt-daemon qemu
qemu-kvm virt-manager
# systemctl enable libvirtd --now
```

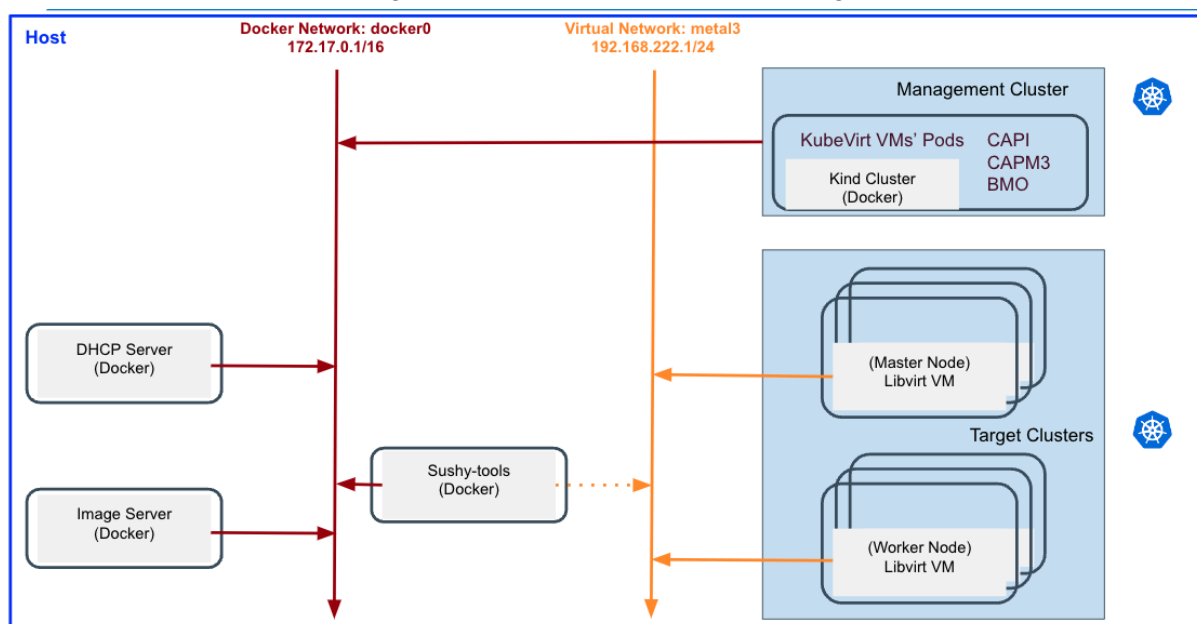
To manage the lifecycle of a Cluster API management cluster, the “clusterctl” CLI tool is needed as well.

```
# curl -L
https://github.com/kubernetes-sigs/cluster-api/releases/download/v1.8.1/clusterctl-
linux-amd64 -o clusterctl
# install -o root -g root -m 0755 clusterctl /usr/local/bin/clusterctl
# clusterctl version
```

What is more, you will need create and update the flat-files used to store usernames and password for basic authentication of Ironi, so install the following package.

```
# apt install apache2-utils
```

To have a better understanding of the environment, here is the figure.



Now, you are 100% ready to go, and it will be a plus if you are in a GOOD MOOD cause there will be, well, some unexpected and bitchy troubles most probably.

1. Set up the bootstrap Kubernetes cluster via [kind](#)

NOTE: This bootstrap Kubernetes cluster will serve as the management Kubernetes cluster later.

Pls, refer to [Installation](#) to install “kind” and choose whichever way you like, in this doc, “kind” is installed via “go install”. After the installation, the “kind” command should be available.






```
# Download the go package from https://go.dev/doc/install

# tar -C /usr/local -xzf go1.22.5.linux-amd64.tar.gz
# echo "export GOPATH=/root/go" >> ~/.bashrc
# echo "export GOROOT=/usr/local/go" >> ~/.bashrc
# echo "export GO111MODULE=on" >> ~/.bashrc
# echo "export GOPROXY=https://goproxy.cn,direct" >> ~/.bashrc
# echo "export PATH=$PATH:/usr/local/go/bin:/root/go/bin" >> ~/.bashrc
```

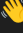
```
# source ~/.bashrc
# go install sigs.k8s.io/kind@latest
# kind --version
kind version 0.24.0
```

Then follow the following commands to create the Kubernetes cluster with the specific configuration. At the end, you should be able to see a Kubernetes cluster is up, and it is running inside a docker container (which is really cool!!!).

```
# cat kind.yaml
kind: Cluster
apiVersion: kind.x-k8s.io/v1alpha4
nodes:
- role: control-plane
  # Open ports for IroniC
  extraPortMappings:
  # IroniC httpd
  - containerPort: 6180
    hostPort: 6180
    listenAddress: "0.0.0.0"
    protocol: TCP
  # IroniC API
  - containerPort: 6385
    hostPort: 6385
    listenAddress: "0.0.0.0"
    protocol: TCP
  # Inspector API
  - containerPort: 5050
    hostPort: 5050
    listenAddress: "0.0.0.0"
    protocol: TCP
  extraMounts:
  - containerPath: /var/lib/kubelet/config.json
    hostPath: /etc/docker/daemon.json

# kind create cluster --config kind.yaml
Creating cluster "kind" ...
✓ Ensuring node image (kindest/node:v1.31.0) 
✓ Preparing nodes 
✓ Writing configuration 
✓ Starting control-plane 
✓ Installing StorageClass 
Set kubectl context to "kind-kind"
You can now use your cluster with:

kubectl cluster-info --context kind-kind

Have a nice day! 

# kubectl cluster-info --context kind-kind
Kubernetes control plane is running at https://127.0.0.1:42937
CoreDNS is running at
https://127.0.0.1:42937/api/v1/namespaces/kube-system/services/kube-dns:dns/proxy

To further debug and diagnose cluster problems, use 'kubectl cluster-info dump'.

# docker container ps -a
CONTAINER ID   IMAGE                                COMMAND                                  CREATED
STATUS        PORTS
NAMES
f7d23f56943c   kindest/node:v1.31.0               "/usr/local/bin/entr..."             About a minute ago
Up About a minute   0.0.0.0:5050->5050/tcp, 0.0.0.0:6180->6180/tcp,
0.0.0.0:6385->6385/tcp, 127.0.0.1:42937->6443/tcp   kind-control-plane
```

You could simply check the pods in the Kubernetes cluster with “kubectl” commands (“coredns” and “local-path-storage” are not ready because there is no network plugin deployed yet).

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
cert-manager	cert-manager-7fbbc65b49-khr9p	0/1	Pending	0	14s
cert-manager	cert-manager-cainjector-6664fc84f6-gkwmg	0/1	Pending	0	14s
cert-manager	cert-manager-webhook-59598898fd-sxf9w	0/1	Pending	0	14s
kube-system	coredns-6f6b679f8f-9rf8t	0/1	Pending	0	70m
kube-system	coredns-6f6b679f8f-fg7fg	0/1	Pending	0	70m
kube-system	etcd-kind-control-plane	1/1	Running	0	70m
kube-system	kube-apiserver-kind-control-plane	1/1	Running	0	70m
kube-system	kube-controller-manager-kind-control-plane	1/1	Running	0	70m
kube-system	kube-proxy-5kqms	1/1	Running	0	70m
kube-system	kube-scheduler-kind-control-plane	1/1	Running	0	70m
local-path-storage	local-path-provisioner-57c5987fd4-hbcpg	0/1	Pending	0	70m

2. Deploy Calico CNI

Pls, refer to [installation](#) to install “Calico CNI” and choose whichever way you like, in this doc, “Calico CNI” is installed by simply using the Calico manifest to create the required resources. After the installation, you should be able to see Calico-related deployments and pods are up. Once the network plugin is ready, “coredns” and “local-path-storage” should be ready too.

```
# wget
https://raw.githubusercontent.com/projectcalico/calico/master/manifests/calico.yaml
# kubectl apply -f calico.yaml
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
kube-system	calico-kube-controllers-7fbd86d5c5-4drcg	1/1	Running	0	10m
kube-system	calico-node-zc4rb	1/1	Running	0	9m45s
kube-system	coredns-6f6b679f8f-9rf8t	1/1	Running	0	83m
kube-system	coredns-6f6b679f8f-fg7fg	1/1	Running	0	83m
kube-system	etcd-kind-control-plane	1/1	Running	0	83m
kube-system	kube-apiserver-kind-control-plane	1/1	Running	0	83m
kube-system	kube-controller-manager-kind-control-plane	1/1	Running	0	83m
kube-system	kube-proxy-5kqms	1/1	Running	0	83m
kube-system	kube-scheduler-kind-control-plane	1/1	Running	0	83m
local-path-storage	local-path-provisioner-57c5987fd4-hbcpg	1/1	Running	0	83m

3. Deploy Cert Manager

```
# wget
https://github.com/cert-manager/cert-manager/releases/download/v1.15.3/cert-manager.yaml
# kubectl apply -f cert-manager.yaml
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -n cert-manager
NAME                                READY   STATUS    RESTARTS   AGE
cert-manager-7fbbc65b49-vjvpb      1/1     Running   0           16s
cert-manager-cainjector-6664fc84f6-lpcxh 1/1     Running   0           16s
cert-manager-webhook-59598898fd-cftb8 1/1     Running   0           16s
```

4. Start VMs

NOTE: As I do not have any bare metal machines yet, so I prepared the VMs to simulate bare metal machines.

First, prepare a virtual network for the VMs.

```
# cat baremetal.xml
<network>
  <name>baremetal</name>
  <forward mode='nat'>
    <nat>
      <port start='1024' end='65535' />
    </nat>
  </forward>
  <bridge name='metal3' />
  <ip address='192.168.222.1' netmask='255.255.255.0'>
  </ip>
</network>

# virsh net-define baremetal.xml
# virsh net-start baremetal
# virsh net-autostart baremetal
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# virsh net-list
Name          State    Autostart  Persistent
-----
baremetal     active  yes        yes
default       active  yes        yes
```

Use “virt-install” to start two VMs (At this moment, VMs do not have any OS running inside).

```
# virt-install \
--connect qemu:///system \
--name bmh-vm-01 \
--description "Virtualized BareMetalHost" \
--osinfo=ubuntu-lts-latest \
--ram=8192 \
--vcpus=2 \
--disk size=25 \
--graphics=none \
--console pty \
--serial pty \
--pxe \
--network network=baremetal,mac="00:60:2f:31:81:01" \
--noautoconsole

virt-install \
--connect qemu:///system \
--name bmh-vm-02 \
--description "Virtualized BareMetalHost" \
--osinfo=ubuntu-lts-latest \
--ram=8192 \
```

```
--vcpus=2 \
--disk size=25 \
--graphics=none \
--console pty \
--serial pty \
--pxe \
--network network=baremetal,mac="00:60:2f:31:81:02" \
--noautoconsole
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# virsh list --all
 Id   Name          State
-----
 13   bmh-vm-01     running
 14   bmh-vm-02     running
```

5. Set up sushy-tools to manage VMs

Metal3 relies on BMC to manage the bare metal machines, so we need something similar for the VMs. This comes in the form of sushy-tools.

```
# cat sushy-tools.conf
# Listen on 192.168.222.1:8000
SUSHY_EMULATOR_LISTEN_IP = u'192.168.222.1'
SUSHY_EMULATOR_LISTEN_PORT = 8000
# The libvirt URI to use. This option enables libvirt driver.
SUSHY_EMULATOR_LIBVIRT_URI = u'qemu:///system'

# docker run --name sushy-tools --rm --network host -d \
-v /var/run/libvirt:/var/run/libvirt \
-v "$(pwd)/sushy-tools.conf:/etc/sushy/sushy-emulator.conf" \
-e SUSHY_EMULATOR_CONFIG=/etc/sushy/sushy-emulator.conf \
quay.io/metal3-io/sushy-tools:latest sushy-emulator
```

Via the Redfish REST API, you should be able to use GET method to check the VMs information and use POST method to power on/off the VMs.

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# virsh domuuid bmh-vm-01
341c614e-46dd-4c34-b6cc-5312c15cf29c

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# virsh domuuid bmh-vm-02
5778cdde-10dd-4dcc-b58d-c551d32c6d5f

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# curl -s http://192.168.222.1:8000/redfish/v1/Systems | jq '.Members'
[
  {
    "@odata.id": "/redfish/v1/Systems/5778cdde-10dd-4dcc-b58d-c551d32c6d5f"
  },
  {
    "@odata.id": "/redfish/v1/Systems/341c614e-46dd-4c34-b6cc-5312c15cf29c"
  }
]
```



```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# curl -s http://192.168.222.1:8000/redfish/v1/Systems/341c614e-46dd-4c34-b6cc-5312c15cf29c | jq '.Actions'
{
  "#ComputerSystem.Reset": {
    "target": "/redfish/v1/Systems/341c614e-46dd-4c34-b6cc-5312c15cf29c/Actions/ComputerSystem.Reset",
    "ResetType@Redfish.AllowableValues": [
      "On",
      "ForceOff",
      "GracefulShutdown",
      "GracefulRestart",
      "ForceRestart",
      "Nmi",
      "ForceOn"
    ]
  }
}

```

6. Set up DHCP Server for VMs

In reality, PXE is usually used when provisioning bare metal machines, a DHCP server is needed to assign the bare metal machines proper IPs.

```

# cat dnsmasq.env
DHCP_HOSTS=00:60:2f:31:81:01;00:60:2f:31:81:02
DHCP_IGNORE=tag:!known
# IP of the host from VM perspective
PROVISIONING_IP=192.168.222.1
GATEWAY_IP=192.168.222.1
DHCP_RANGE=192.168.222.100,192.168.222.149

# docker run --name dnsmasq --rm -d --net=host --privileged --user 997:994 \
--env-file dnsmasq.env --entrypoint /bin/rundnsmasq \
quay.io/metal3-io/ironic

```

7. Set up Image Server for VMs

In reality, when provisioning bare metal machines, iso files will be pulled to install the OS via PXE for the bare metal machines. Thus, an Image server is needed.

```

# mkdir disk-images
# cd disk-images
#
# https://cloud-images.ubuntu.com/jammy/current/jammy-server-cloudimg-amd64.img wget
# wget https://cloud-images.ubuntu.com/jammy/current/SHA256SUMS
# sha256sum --ignore-missing -c SHA256SUMS
#
# https://cloud.centos.org/centos/9-stream/x86_64/images/CentOS-Stream-GenericCloud-9-latest.x86_64.qcow2 wget
#
# https://cloud.centos.org/centos/9-stream/x86_64/images/CentOS-Stream-GenericCloud-9-latest.x86_64.qcow2.SHA256SUM wget
# sha256sum -c CentOS-Stream-GenericCloud-9-latest.x86_64.qcow2.SHA256SUM
#
# https://artifactory.nordix.org/artifactory/metal3/images/k8s_v1.29.0/CENTOS_9_NODE_IMAGE_K8S_v1.29.0.qcow2 wget
# sha256sum CENTOS_9_NODE_IMAGE_K8S_v1.29.0.qcow2
# cd ..
# docker run --name image-server --rm -d -p 80:8080 \
-v "$(pwd)/disk-images:/usr/share/nginx/html" nginxinc/nginx-unprivileged

```

```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# curl --head http://127.0.0.1:80/CENTOS_9_NODE_IMAGE_K8S_v1.29.0.qcow2
HTTP/1.1 200 OK
Server: nginx/1.27.1
Date: Mon, 19 Aug 2024 08:43:40 GMT
Content-Type: application/octet-stream
Content-Length: 2270668288
Last-Modified: Wed, 10 Jan 2024 12:51:34 GMT
Connection: keep-alive
ETag: "659e92d6-8757a600"
Accept-Ranges: bytes

```

8. Deploy Ironic for Metal³

```

# tree ironic/
ironic/
├── ironic-auth-config
├── ironic_bmo.env
├── ironic-htpasswd
├── ironic-inspector-auth-config
├── ironic-inspector-htpasswd
├── ironic-patch.yaml
└── kustomization.yaml

# cat ironic-auth-config
[ironic]
auth_type=http_basic
username=IRONIC_USERNAME
password=IRONIC_PASSWORD

# cat ironic-inspector-auth-config
[inspector]
auth_type=http_basic
username=INSPECTOR_USERNAME
password=INSPECTOR_PASSWORD

# cat ironic-htpasswd
IRONIC_HTPASSWD=<output of `htpasswd -n -b -B IRONIC_USERNAME IRONIC_PASSWORD`>

# cat ironic-inspector-htpasswd
INSPECTOR_HTPASSWD=<output of `htpasswd -n -b -B INSPECTOR_USERNAME
INSPECTOR_PASSWORD`>

# cat ironic_bmo.env
HTTP_PORT=6180
PROVISIONING_INTERFACE=eth0
CACHEURL=http://192.168.222.1/images
IRONIC_KERNEL_PARAMS=console=ttyS0

# cat ironic-patch.yaml
apiVersion: apps/v1
kind: Deployment
metadata:
  name: ironic
spec:
  template:
    spec:
      containers:
        - name: ironic-dnsmasq
          $patch: delete

# cat kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1

```

```

kind: Kustomization
namespace: baremetal-operator-system
resources:
- ../baremetal-operator-main/config/namespace
- ../baremetal-operator-main/ironic-deployment/base
# The kustomize components configure basic-auth and TLS
components:
- ../baremetal-operator-main/ironic-deployment/components/basic-auth
- ../baremetal-operator-main/ironic-deployment/components/tls
images:
- name: quay.io/metal3-io/ironic
  newTag: latest
# Create a ConfigMap from ironic_bmo.env and call it ironic-bmo-configmap.
# This ConfigMap will be used to set environment variables for the containers.
configMapGenerator:
- envs:
  - ironic_bmo.env
  name: ironic-bmo-configmap
  behavior: create

patches:
# Patch for removing dnsmasq
- path: ironic-patch.yaml
# The TLS component adds certificates but it cannot know the exact IPs of our
environment.
# Here we patch the certificates to have the correct IPs.
# - 172.18.0.2: kind cluster node IP. This is what Ironic will see attached to the
interface
# and use to communicate with Inspector.
- patch: |-
  - op: replace
    path: /spec/ipAddresses/0
    value: 192.168.222.1
  - op: add
    path: /spec/ipAddresses/-
    value: 172.18.0.2
  target:
    kind: Certificate
    name: ironic-cert|ironic-inspector-cert
# The CA certificate should not have any IP address so we remove it.
- patch: |-
  - op: remove
    path: /spec/ipAddresses
  target:
    kind: Certificate
    name: ironic-cacert
# Create secrets from the authentication configuration.
# These will be mounted or used for environment variables.
# See the basic-auth component for more details on how they are used.
secretGenerator:
- name: ironic-htpasswd
  behavior: create
  envs:
  - ironic-htpasswd
- name: ironic-inspector-htpasswd
  behavior: create
  envs:
  - ironic-inspector-htpasswd
- name: ironic-auth-config
  files:
  - auth-config=ironic-auth-config
- name: ironic-inspector-auth-config
  files:
  - auth-config=ironic-inspector-auth-config

# kubectl create -k ironic --dry-run=client -o yaml
# kubectl apply -k ironic

```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get deployment ironic -n baremetal-operator-system
NAME      READY   UP-TO-DATE   AVAILABLE   AGE
ironic    1/1     1             1           81m
```

9. Deploy Bare Metal Operator for Metal³

```
# tree bmo/
bmo/
├── ironic.env
├── ironic-inspector-password
├── ironic-inspector-username
├── ironic-password
├── ironic-username
└── kustomization.yaml

# cat ironic-username
IRONIC_USERNAME

# cat ironic-password
IRONIC_PASSWORD

# cat ironic-inspector-username
INSPECTOR_USERNAME

# cat ironic-inspector-password
INSPECTOR_PASSWORD

# cat ironic.env
DEPLOY_KERNEL_URL=http://192.168.222.1:6180/images/ironic-python-agent.kernel
DEPLOY_RAMDISK_URL=http://192.168.222.1:6180/images/ironic-python-agent.initramfs
IRONIC_ENDPOINT=https://192.168.222.1:6385/v1/

# cat kustomization.yaml
apiVersion: kustomize.config.k8s.io/v1beta1
kind: Kustomization
namespace: baremetal-operator-system
resources:
- ../baremetal-operator-main/config/overlays/basic-auth_tls
images:
- name: quay.io/metal3-io/baremetal-operator
  newTag: latest
# Create a ConfigMap from ironic.env and name it ironic.
configMapGenerator:
- name: ironic
  behavior: create
  envs:
  - ironic.env

# We cannot use suffix hashes since the kustomizations we build on
# cannot be aware of what suffixes we add.
generatorOptions:
  disableNameSuffixHash: true
# Create secrets with the credentials for accessing Ironic.
secretGenerator:
- name: ironic-credentials
  files:
  - username=ironic-username
  - password=ironic-password
- name: ironic-inspector-credentials
  files:
  - username=ironic-inspector-username
  - password=ironic-inspector-password
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get deployment baremetal-operator-controller-manager
-n baremetal-operator-system
```

NAME	READY	UP-TO-DATE	AVAILABLE	AGE
baremetal-operator-controller-manager	1/1	1	1	19m

10. Create BareMetalHosts

```
# cat bml-vm-01.yaml
apiVersion: v1
kind: Secret
metadata:
  name: bml-01
type: Opaque
stringData:
  username: replaceme
  password: replaceme
---
apiVersion: metal3.io/v1alpha1
kind: BareMetalHost
metadata:
  name: bml-vm-01
spec:
  online: true
  bootMACAddress: 00:60:2f:31:81:01
  bootMode: UEFI # use 'legacy' for Scenario 2
  hardwareProfile: libvirt
  bmc:
    address:
redfish-virtualmedia+http://192.168.222.1:8000/redfish/v1/Systems/341c614e-46dd-4c3
4-b6cc-5312c15cf29c
    credentialsName: bml-01

# cat bml-vm-02.yaml
apiVersion: v1
kind: Secret
metadata:
  name: bml-02
type: Opaque
stringData:
  username: replaceme
  password: replaceme
---
apiVersion: metal3.io/v1alpha1
kind: BareMetalHost
metadata:
  name: bml-vm-02
spec:
  online: true
  bootMACAddress: 00:60:2f:31:81:02
  bootMode: UEFI # use 'legacy' for Scenario 2
  hardwareProfile: libvirt
  bmc:
    address:
redfish-virtualmedia+http://192.168.222.1:8000/redfish/v1/Systems/5778cdde-10dd-4dc
c-b58d-c551d32c6d5f
    credentialsName: bml-02

# kubectl apply -f bml-vm-01.yaml
# kubectl apply -f bml-vm-02.yaml
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get bmh -A
NAMESPACE   NAME           STATE      CONSUMER  ONLINE  ERROR  AGE
default     bml-vm-01     available
default     bml-vm-02     available
              true     9m52s
              true     8m33s
```

11. Install MetalLB

KubeVirt is a cloud native virtualization solution. The VMs we're going to create and use for the workload cluster's nodes, are actually running within pods in the management cluster. In order to communicate with the workload cluster's API server, we'll need to expose it. The easiest way to expose the workload cluster's API server (a pod within a node running in a VM that is itself running within a pod in the management cluster, that is running inside a Docker container), is to use a LoadBalancer service, namely MetalLB in this doc.

```
# wget
https://raw.githubusercontent.com/metallb/metallb/v0.14.8/config/manifests/metallb-
native.yaml
# kubectl apply -f metallb-native.yaml

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get deployments -n metallb-system
NAME          READY  UP-TO-DATE  AVAILABLE  AGE
controller    1/1    1           1          8m15s
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~#
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -n metallb-system
NAME                                READY  STATUS    RESTARTS  AGE
controller-8694df9d9b-jgr57        1/1    Running   0          8m29s
speaker-7cjcfc                      1/1    Running   0          8m29s
```

Create the IPAddressPool and the L2Advertisement custom resources.

```
# cat capi-ip-pool.yaml
apiVersion: metallb.io/v1beta1
kind: IPAddressPool
metadata:
  name: capi-ip-pool
  namespace: metallb-system
spec:
  addresses:
  - 172.18.255.200-172.18.255.250
---
apiVersion: metallb.io/v1beta1
kind: L2Advertisement
metadata:
  name: empty
  namespace: metallb-system
# kubectl apply -f capi-ip-pool.yaml
```

12. Deploy KubeVirt

```
# wget
https://github.com/kubevirt/kubevirt/releases/download/v1.3.0/kubevirt-operator.yam
l
# wget
https://github.com/kubevirt/kubevirt/releases/download/v1.3.0/kubevirt-cr.yaml
# kubectl apply -f kubevirt-operator.yaml
```

```
# kubectl apply -f kubevirt-cr.yaml
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -n kubevirt
NAME                                READY   STATUS    RESTARTS   AGE
virt-api-79f4646554-h7mhp           1/1     Running   0           11m
virt-controller-7cbbdbbf8f-pjmcr    1/1     Running   0           10m
virt-controller-7cbbdbbf8f-v6jtj    1/1     Running   0           10m
virt-handler-z2dwv                  1/1     Running   0           10m
virt-operator-84d89fd9f6-cj9zl      1/1     Running   0           13m
virt-operator-84d89fd9f6-rbjd6      1/1     Running   0           13m
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~#
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get deployments -n kubevirt
NAME                READY   UP-TO-DATE   AVAILABLE   AGE
virt-api            1/1     1             1           12m
virt-controller     2/2     2             2           10m
virt-operator       2/2     2             2           13m
```

13. Initialize the management cluster

In this chapter, steps to initialize the management cluster with the KubeVirt Provider and Metal³ Provider will be displayed.

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# clusterctl init --infrastructure kubevirt --infrastructure metal3

Fetching providers
Skipping installing cert-manager as it is already installed
Installing provider="cluster-api" version="v1.8.1" targetNamespace="capi-system"
Installing provider="bootstrap-kubeadm" version="v1.8.1" targetNamespace="capi-kubeadm-bootstrap-system"
Installing provider="control-plane-kubeadm" version="v1.8.1" targetNamespace="capi-kubeadm-control-plane-system"
Installing provider="infrastructure-kubevirt" version="v0.1.9" targetNamespace="capk-system"
Installing provider="infrastructure-metal3" version="v1.7.1" targetNamespace="capm3-system"

Your management cluster has been initialized successfully!

You can now create your first workload cluster by running the following:

clusterctl generate cluster [name] --kubernetes-version [version] | kubectl apply -f -
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pod -n capi-system
NAME                                READY   STATUS    RESTARTS   AGE
capi-controller-manager-6cb7846fdf-672ms  1/1     Running   0           56m
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~#
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pod -n capm3-system
NAME                                READY   STATUS    RESTARTS   AGE
capm3-controller-manager-864b9cddc6-4t5l9  1/1     Running   0           56m
ipam-controller-manager-64f4697b6b-6pnrp  1/1     Running   0           56m
```

After all of these, you should be able to get a management Kubernetes cluster with the following pods/deployments/services.


```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -A
```

NAMESPACE	NAME	READY	STATUS	RESTARTS	AGE
baremetal-operator-system	baremetal-operator-controller-manager-8d59474bd-ndmg7	2/2	Running	0	39h
baremetal-operator-system	ironic-8569ccdcfb-mdkpw	3/3	Running	0	40h
capi-kubeadm-bootstrap-system	capi-kubeadm-bootstrap-controller-manager-7989ff8dd5-fs7jr	1/1	Running	0	38h
capi-kubeadm-control-plane-system	capi-kubeadm-control-plane-controller-manager-5df44cc55f-cml8s	1/1	Running	0	38h
capi-system	capi-controller-manager-6cb7846fdf-672ms	1/1	Running	0	38h
capk-system	capk-controller-manager-777f8c496b-9shzg	1/1	Running	0	38h
capm3-system	capm3-controller-manager-864b9cddc6-4t5l9	1/1	Running	0	38h
capm3-system	ipam-controller-manager-64f4697b6b-6pnrp	1/1	Running	0	38h
cert-manager	cert-manager-7fbbc65b49-vjvpb	1/1	Running	0	45h
cert-manager	cert-manager-cainjector-6664fc84f6-lpcxh	1/1	Running	0	45h
cert-manager	cert-manager-webhook-59598898fd-cftb8	1/1	Running	0	45h
default	virt-launcher-damn-kubevirt-control-plane-zmdq9-nvnsw	3/3	Running	0	19h
default	virt-launcher-damn-kubevirt-md-0-6j2sz-b2drk-nlwg7	3/3	Running	0	19h
default	virt-launcher-mixed-quickstart-control-plane-rgz2b-8zrjj	3/3	Running	0	19h
kube-system	calico-kube-controllers-7fbd86d5c5-4drcg	1/1	Running	0	45h
kube-system	calico-node-zc4rb	1/1	Running	0	45h
kube-system	coredns-6f6b679f8f-9rf8t	1/1	Running	0	47h
kube-system	coredns-6f6b679f8f-fg7fg	1/1	Running	0	47h
kube-system	etcd-kind-control-plane	1/1	Running	0	47h
kube-system	kube-apiserver-kind-control-plane	1/1	Running	0	47h
kube-system	kube-controller-manager-kind-control-plane	1/1	Running	0	47h
kube-system	kube-proxy-5kqms	1/1	Running	0	47h
kube-system	kube-scheduler-kind-control-plane	1/1	Running	0	47h
kubevirt	virt-api-79f4646554-h7mhp	1/1	Running	0	39h
kubevirt	virt-controller-7cbbdbbf8f-pjmcr	1/1	Running	0	38h
kubevirt	virt-controller-7cbbdbbf8f-v6jttj	1/1	Running	0	38h
kubevirt	virt-handler-z2dwv	1/1	Running	0	38h
kubevirt	virt-operator-84d89fd9f6-cj9zl	1/1	Running	0	39h
kubevirt	virt-operator-84d89fd9f6-rbjd6	1/1	Running	0	39h
local-path-storage	local-path-provisioner-57c5987fd4-hbcpd	1/1	Running	0	47h
metallb-system	controller-8694df9d9b-jgr57	1/1	Running	0	39h
metallb-system	speaker-7cjcj	1/1	Running	0	39h

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get deployments -A
```

NAMESPACE	NAME	READY	UP-TO-DATE	AVAILABLE	AGE
baremetal-operator-system	baremetal-operator-controller-manager	1/1	1	1	39h
baremetal-operator-system	ironic	1/1	1	1	41h
capi-kubeadm-bootstrap-system	capi-kubeadm-bootstrap-controller-manager	1/1	1	1	38h
capi-kubeadm-control-plane-system	capi-kubeadm-control-plane-controller-manager	1/1	1	1	38h
capi-system	capi-controller-manager	1/1	1	1	38h
capk-system	capk-controller-manager	1/1	1	1	38h
capm3-system	capm3-controller-manager	1/1	1	1	38h
capm3-system	ipam-controller-manager	1/1	1	1	38h
cert-manager	cert-manager	1/1	1	1	45h
cert-manager	cert-manager-cainjector	1/1	1	1	45h
cert-manager	cert-manager-webhook	1/1	1	1	45h
kube-system	calico-kube-controllers	1/1	1	1	45h
kube-system	coredns	2/2	2	2	47h
kubevirt	virt-api	1/1	1	1	39h
kubevirt	virt-controller	2/2	2	2	38h
kubevirt	virt-operator	2/2	2	2	39h
local-path-storage	local-path-provisioner	1/1	1	1	47h
metallb-system	controller	1/1	1	1	39h

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get services -A
```

NAMESPACE	NAME	TYPE	CLUSTER-IP	EXTERNAL-IP	PORT(S)	AGE
baremetal-operator-system	baremetal-operator-controller-manager-metrics-service	ClusterIP	10.96.197.130	<none>	8443/TCP	39h
baremetal-operator-system	baremetal-operator-webhook-service	ClusterIP	10.96.62.32	<none>	443/TCP	39h
capi-kubeadm-bootstrap-system	capi-kubeadm-bootstrap-webhook-service	ClusterIP	10.96.47.228	<none>	443/TCP	38h
capi-kubeadm-control-plane-system	capi-kubeadm-control-plane-webhook-service	ClusterIP	10.96.156.3	<none>	443/TCP	38h
capi-system	capi-webhook-service	ClusterIP	10.96.35.19	<none>	443/TCP	38h
capk-system	capk-webhook-service	ClusterIP	10.96.102.66	<none>	443/TCP	38h
capm3-system	capm3-webhook-service	ClusterIP	10.96.32.140	<none>	443/TCP	38h
capm3-system	ipam-webhook-service	ClusterIP	10.96.102.170	<none>	443/TCP	38h
cert-manager	cert-manager	ClusterIP	10.96.138.108	<none>	9402/TCP	45h
cert-manager	cert-manager-webhook	ClusterIP	10.96.116.20	<none>	443/TCP	45h
default	damn-kubevirt-lb	LoadBalancer	10.96.46.29	172.18.0.200	6443:30739/TCP	19h
default	kubernetes	ClusterIP	10.96.0.1	<none>	443/TCP	47h
kube-system	kube-dns	ClusterIP	10.96.0.10	<none>	53/UDP,53/TCP,9153/TCP	47h
kubevirt	kubevirt-operator-webhook	ClusterIP	10.96.227.49	<none>	443/TCP	39h
kubevirt	kubevirt-prometheus-metrics	ClusterIP	None	<none>	443/TCP	39h
kubevirt	virt-api	ClusterIP	10.96.220.50	<none>	443/TCP	39h
kubevirt	virt-exportproxy	ClusterIP	10.96.205.118	<none>	443/TCP	39h
metallb-system	metallb-webhook-service	ClusterIP	10.96.15.38	<none>	443/TCP	39h

14. Create the target cluster on KubeVirt VMs only

```
# export CAPK_GUEST_K8S_VERSION="v1.30.1"
# export CRI_PATH="/var/run/containerd/containerd.sock"
# export NODE_VM_IMAGE_TEMPLATE="quay.io/capik/centos-2204-container-disk:v1.30.1"
```



```
# clusterctl generate cluster damn-kubevirt \
--infrastructure="kubevirt" \
--flavor lb \
--kubernetes-version ${CAPK_GUEST_K8S_VERSION} \
--control-plane-machine-count=1 \
--worker-machine-count=1 \
> damn-kubevirt.yaml
# kubectl apply -f damn-kubevirt.yaml
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get cluster
NAME                CLUSTERCLASS  PHASE      AGE      VERSION
damn-kubevirt              Provisioned   3m35s
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get kubevirtmachines
NAME                                AGE      READY
damn-kubevirt-control-plane-zmdq9  3m36s    true
damn-kubevirt-md-0-6j2sz-b2drk     3m36s    true
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get pods -A | grep damn
default          virt-launcher-damn-kubevirt-control-plane-zmdq9-nvsw      3/3      Running    0          12m
default          virt-launcher-damn-kubevirt-md-0-6j2sz-b2drk-nlzg7        3/3      Running    0          11m
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get machines -A
NAMESPACE  NAME                                CLUSTER      NODENAME
PROVIDERID
default    damn-kubevirt-control-plane-zmdq9  damn-kubevirt  damn-kubevirt-control-plane-zmdq9
kubevirt://damn-kubevirt-control-plane-zmdq9  Running  3m41s  v1.30.1
default    damn-kubevirt-md-0-6j2sz-b2drk     damn-kubevirt  damn-kubevirt-md-0-6j2sz-b2drk
kubevirt://damn-kubevirt-md-0-6j2sz-b2drk     Running  3m41s  v1.30.1
```

```
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# clusterctl describe cluster damn-kubevirt
NAME                                READY  SEVERITY  REASON  SINCE  MESSAGE
Cluster/damn-kubevirt
├─ClusterInfrastructure - KubevirtCluster/damn-kubevirt      True      10m
├─ControlPlane - KubeadmControlPlane/damn-kubevirt-control-plane  True      11m
├─Machine/damn-kubevirt-control-plane-zmdq9      True      10m
└─Workers
  └─MachineDeployment/damn-kubevirt-md-0          False  Warning  WaitingForAvailableMachines  11m  Minimum a
    Machine/damn-kubevirt-md-0-6j2sz-b2drk      True      10m
```

15. Create the target cluster on Bare Metal Machines only (**problematic**)

```
# export
IMAGE_CHECKSUM="7b2fbe69b2f2446d151b3e198b7fb020a4f17b9ab237c1b59b843e2783218b66"
# export IMAGE_CHECKSUM_TYPE="sha256"
# export IMAGE_FORMAT="qcow2"
# export IMAGE_URL="http://192.168.222.1/UBUNTU_22.04_NODE_IMAGE_K8S_v1.30.0.qcow2"
# export KUBERNETES_VERSION="v1.30.0"
# export POD_CIDR='["192.168.10.0/24"]'
# export CTLPLANE_KUBEADM_EXTRA_CONFIG=""
# export WORKERS_KUBEADM_EXTRA_CONFIG=""
# export CLUSTER_APIENDPOINT_HOST="192.168.222.100"
# export CLUSTER_APIENDPOINT_PORT="6443"
# clusterctl generate cluster stupid-metal3 --control-plane-machine-count 1
--worker-machine-count 1 --infrastructure="metal3" > fuck-metal3.yaml
# kubectl apply -f fuck-metal3.yaml
```

```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get cluster fuck-metal3
NAME          CLUSTERCLASS  PHASE      AGE    VERSION
fuck-metal3    Provisioned   141m
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~#
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl get bmh -A
NAMESPACE  NAME          STATE      CONSUMER          ONLINE  ERROR  AGE
default    bml-vm-01     provisioned fuck-metal3-dtkg2  true    144m

```

Indeed, there is an issue left without solving. As we can see from the pictures, the cluster is created and the “bmh” resource is consumed by the cluster successfully. What is more, the corresponding VM for the “bmh” resource is installed with OS via PXE and obtained the IP address. From the laptop, “ping” this “bmh” resource also succeeded.

```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# kubectl describe bmh bml-vm-01
Name:          bml-vm-01
Namespace:     default
Labels:        cluster.x-k8s.io/cluster-name=fuck-metal3
Annotations:   <none>
API Version:   metal3.io/v1alpha1
Kind:          BareMetalHost
Metadata:
  Creation Timestamp:  2024-08-21T03:27:31Z
  Finalizers:
    baremetalhost.metal3.io
  Generation:  4
  Owner References:
    API Version:  infrastructure.cluster.x-k8s.io/v1beta1
    Controller:   true
    Kind:         Metal3Machine
    Name:         fuck-metal3-dtkg2
    UID:          90a82087-b5b4-4e08-8ddb-40c05152b303
  Resource Version:  714436
  UID:              f2990f89-ba55-4092-92fd-e57f36362cb3
Spec:
  Architecture:      x86_64
  Automated Cleaning Mode:  metadata
  Bmc:
    Address:
redfish-virtualmedia+http://192.168.222.1:8000/redfish/v1/Systems/341c614e-46dd-4c3
4-b6cc-5312c15cf29c
    Credentials Name:  bml-01
  Boot MAC Address:   00:60:2f:31:81:01
  Boot Mode:          UEFI
  Consumer Ref:
    API Version:  infrastructure.cluster.x-k8s.io/v1beta1
    Kind:         Metal3Machine
    Name:         fuck-metal3-dtkg2
    Namespace:    default
  Hardware Profile:   libvirt
  Image:
    Checksum:        7b2fbe69b2f2446d151b3e198b7fb020a4f17b9ab237c1b59b843e2783218b66
    Checksum Type:   sha256
    Format:          qcow2
    URL:             http://192.168.222.1/UBUNTU_22.04_NODE_IMAGE_K8S_v1.30.0.qcow2
  Online:            true
  User Data:
    Name:            fuck-metal3-dtkg2
    Namespace:       default
Status:
  Error Count:       0
  Error Message:
  Good Credentials:
    Credentials:
      Name:          bml-01

```

```
Namespace:      default
Credentials Version: 714034
Hardware:
  Cpu:
    Arch:      x86_64
    Count:     2
    Flags:
      3dnowprefetch
      abm
      adx
      aes
      apic
      arat
      arch_capabilities
      avx
      avx2
      avx_vnni
      bmi1
      bmi2
      clflush
      clflushopt
      clwb
      cmov
      constant_tsc
      cpuid
      cpuid_fault
      cx16
      cx8
      de
      ept
      ept_ad
      erms
      f16c
      flexpriority
      fma
      fpu
      fsgsbase
      fsrm
      fxsr
      gfni
      hypervisor
      ibpb
      ibrs
      ibrs_enhanced
      invpcid
      lahf_lm
      lm
      mca
      mce
      md_clear
      mmx
      movbe
      movdir64b
      movdiri
      msr
      mtrr
      nopl
      nx
      ospke
      pae
      pat
      pclmulqdq
      pdpe1gb
      pge
      pku
      pni
      popcnt
      pse
```

```
pse36
rdpid
rdrand
rdseed
rdtscp
rep_good
sep
serialize
sha_ni
smmap
smep
ss
ssbd
sse
sse2
sse4_1
sse4_2
ssse3
stibp
syscall
tpr_shadow
tsc
tsc_adjust
tsc_deadline_timer
tsc_known_freq
umip
vaes
vme
vmx
vnmi
vpclmulqdq
vpid
waitpkg
x2apic
xgetbv1
xsave
xsaves
xsavesopt
xsaves
xtopology
Model: 13th Gen Intel(R) Core(TM) i7-1360P
Firmware:
Bios:
Date: 02/06/2015
Vendor: EFI Development Kit II / OVMF
Version: 0.0.0
Hostname: localhost.localdomain
Nics:
Ip: 192.168.222.100
Mac: 00:60:2f:31:81:01
Model: 0x1af4 0x0001
Name: enp1s0
Ip: fe80::2f24:badb:8d3:796%enp1s0
Mac: 00:60:2f:31:81:01
Model: 0x1af4 0x0001
Name: enp1s0
Ram Mebibytes: 8192
Storage:
Alternate Names:
/dev/vda
/dev/disk/by-path/pci-0000:04:00.0
Name: /dev/disk/by-path/pci-0000:04:00.0
Rotational: true
Size Bytes: 26843545600
Type: HDD
Vendor: 0x1af4
System Vendor:
Manufacturer: QEMU
```

```

    Product Name: Standard PC (Q35 + ICH9, 2009)
Hardware Profile: libvirt
Last Updated: 2024-08-21T03:32:09Z
Operation History:
  Deprovision:
    End: <nil>
    Start: <nil>
  Inspect:
    End: 2024-08-21T03:29:36Z
    Start: 2024-08-21T03:27:41Z
  Provision:
    End: 2024-08-21T03:32:09Z
    Start: 2024-08-21T03:30:49Z
  Register:
    End: 2024-08-21T03:30:49Z
    Start: 2024-08-21T03:30:49Z
Operational Status: OK
Powered On: true
Provisioning:
  ID: 7af54a37-bc29-4e76-bd34-c279d79e1ba1
  Boot Mode: UEFI
  Image:
    Checksum:
7b2f69b2f2446d151b3e198b7fb020a4f17b9ab237c1b59b843e2783218b66
    Checksum Type: sha256
    Format: qcow2
    URL: http://192.168.222.1/UBUNTU_22.04_NODE_IMAGE_K8S_v1.30.0.qcow2
  Root Device Hints:
    Device Name: /dev/vda
    State: provisioned
Tried Credentials:
  Credentials:
    Name: bml-01
    Namespace: default
    Credentials Version: 714034
Events: <none>

```

```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# ping 192.168.222.100 -c1
PING 192.168.222.100 (192.168.222.100) 56(84) bytes of data.
64 bytes from 192.168.222.100: icmp_seq=1 ttl=64 time=0.204 ms

--- 192.168.222.100 ping statistics ---
1 packets transmitted, 1 received, 0% packet loss, time 0ms
rtt min/avg/max/mdev = 0.204/0.204/0.204/0.000 ms
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~#
root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# virsh console bml-vm-01
Connected to domain 'bml-vm-01'
Escape character is ^] (Ctrl + ])

bml-vm-01 login:

```

However, the “capm3-controller-manager-864b9cddc6-4t5l9” raised the error as follows, which is the root cause that the k8s on VMs (simulating the bare metal machines) failed to be up.

```

# kubectl describe machine fuck-metal3-dtkg2
...
Status:
  Bootstrap Ready: true
Conditions:
  Last Transition Time: 2024-08-21T03:32:12Z

```

```

Message:      1 of 2 completed
Reason:      SettingProviderIDOnNodeFailed
Severity:    Error
Status:      False
Type:        Ready
Last Transition Time: 2024-08-21T03:30:48Z
Status:      True
Type:        BootstrapReady
Last Transition Time: 2024-08-21T03:32:12Z
Message:      error retrieving node, requeuing. Object will be requeued
after 30s
Reason:      SettingProviderIDOnNodeFailed
Severity:    Error
Status:      False
Type:        InfrastructureReady
Last Transition Time: 2024-08-21T03:30:48Z
Reason:      WaitingForNodeRef
Severity:    Info
Status:      False
Type:        NodeHealthy
Last Updated: 2024-08-21T03:30:48Z
Observed Generation: 2
Phase:       Provisioning
Events:      <none>

```

```

# kubectl logs capm3-controller-manager-864b9cddc6-4t519 -n capm3-system
...
I0821 05:59:34.777394      1 metal3machine_manager.go:684] "msg"="Updating
machine" "cluster"="fuck-metal3"
"logger"="controllers.Metal3Machine.Metal3Machine-controller"
"machine"="fuck-metal3-dtkg2" "metal3-cluster"="fuck-metal3"
"metal3-machine"={"Namespace":"default","Name":"fuck-metal3-dtkg2"}
I0821 05:59:34.781713      1 metal3machine_manager.go:1599] "msg"="Metal3data is
ready" "cluster"="fuck-metal3"
"logger"="controllers.Metal3Machine.Metal3Machine-controller"
"machine"="fuck-metal3-dtkg2" "metal3-cluster"="fuck-metal3"
"metal3-machine"={"Namespace":"default","Name":"fuck-metal3-dtkg2"}
I0821 05:59:34.781748      1 metal3machine_manager.go:1120] "msg"="Deleting
nodeReuseLabelName from host, if any" "cluster"="fuck-metal3"
"logger"="controllers.Metal3Machine.Metal3Machine-controller"
"machine"="fuck-metal3-dtkg2" "metal3-cluster"="fuck-metal3"
"metal3-machine"={"Namespace":"default","Name":"fuck-metal3-dtkg2"}
I0821 05:59:34.782435      1 metal3machine_manager.go:729] "msg"="Finished
updating machine" "cluster"="fuck-metal3"
"logger"="controllers.Metal3Machine.Metal3Machine-controller"
"machine"="fuck-metal3-dtkg2" "metal3-cluster"="fuck-metal3"
"metal3-machine"={"Namespace":"default","Name":"fuck-metal3-dtkg2"}
E0821 05:59:34.795936      1 metal3machine_manager.go:1812] "msg"="error while
retrieving nodes" "error"="Get \"https://192.168.222.100:6443/api/v1/nodes\": dial
tcp 192.168.222.100:6443: connect: connection refused" "cluster"="fuck-metal3"
"logger"="controllers.Metal3Machine.Metal3Machine-controller"
"machine"="fuck-metal3-dtkg2" "metal3-cluster"="fuck-metal3"
"metal3-machine"={"Namespace":"default","Name":"fuck-metal3-dtkg2"}
I0821 05:59:34.795962      1 metal3machine_manager.go:1315] "msg"="error
retrieving node, requeuing" "cluster"="fuck-metal3"
"logger"="controllers.Metal3Machine.Metal3Machine-controller"
"machine"="fuck-metal3-dtkg2" "metal3-cluster"="fuck-metal3"
"metal3-machine"={"Namespace":"default","Name":"fuck-metal3-dtkg2"}
E0821 05:59:34.795984      1 metal3machine_controller.go:274] "msg"="Failed to set
the target node providerID" "error"="error retrieving node, requeuing. Object will
be requeued after 30s" "logger"="controllers.Metal3Machine" "providerID"=""
I0821 05:59:50.719091      1 metal3labelsync_controller.go:147] "msg"="Could not
find Node Ref on Machine object, will retry"
"logger"="controllers.Metal3LabelSync.metal3-label-sync-controller"
"metal3-label-sync"={"Namespace":"default","Name":"bml-vm-01"}

```

```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# clusterctl describe cluster fuck-metal3
NAME                                     READY SEVERITY REASON                                     SINCE MESSAGE
Cluster/fuck-metal3                     False Error   SettingProviderIDOnNodeFailed @ Machine/fuck-metal3-dtkg2 150m 0 of 1 completed
├─ClusterInfrastructure - Metal3Cluster/fuck-metal3      True
├─ControlPlane - KubeadmControlPlane/fuck-metal3        False Error   SettingProviderIDOnNodeFailed @ Machine/fuck-metal3-dtkg2 150m 0 of 1 completed
└─Machine/fuck-metal3-dtkg2                      False Error   SettingProviderIDOnNodeFailed                               151m 1 of 2 completed

```

16. Create the target cluster on Bare Metal Machines and KubeVirt VMs (problematic)

Well, as we are still using VMs to simulate the bare metal machines here, so there is no doubt we will hit the issue in the last chapter. But I will provide a way here anyway about how to start a mixed cluster.

```

root@kubevirt-ThinkPad-X1-Carbon-Gen-11:~# clusterctl describe cluster damn-kubevirt
NAME                                     READY SEVERITY REASON
Cluster/damn-kubevirt                  True
├─ClusterInfrastructure - KubevirtCluster/damn-kubevirt      True
├─ControlPlane - KubeadmControlPlane/damn-kubevirt-control-plane True
├─Machine/damn-kubevirt-control-plane-zmdq9                  True
└─Workers
  └─MachineDeployment/damn-kubevirt-md-0                      False Warning Waiting
    └─Machine/damn-kubevirt-md-0-6j2sz-b2drk                  True

```

So follow a similar topology, use “KubeVirt VM” as the control plane and “Metal3 Machines” as the workers, the yaml file should be in the following format.

```

# cat final.yaml
apiVersion: cluster.x-k8s.io/v1beta1
kind: Cluster
metadata:
  name: mixed-quickstart
  namespace: default
spec:
  clusterNetwork:
    pods:
      cidrBlocks:
        - 192.168.0.0/16
    services:
      cidrBlocks:
        - 10.128.0.0/16
  controlPlaneRef:
    apiVersion: controlplane.cluster.x-k8s.io/v1beta1
    kind: KubeadmControlPlane
    name: mixed-quickstart-control-plane
    namespace: default
  infrastructureRef:
    apiVersion: infrastructure.cluster.x-k8s.io/v1alpha1
    kind: KubevirtCluster
    name: mixed-quickstart
    namespace: default
---
apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
kind: Metal3Cluster
metadata:
  name: mixed-quickstart
  namespace: default
spec:
  controlPlaneEndpoint:

```

```

    host: 192.168.0.101
    port: 6443
    noCloudProvider: true
---
apiVersion: cluster.x-k8s.io/v1beta1
kind: MachineDeployment
metadata:
  labels:
    cluster.x-k8s.io/cluster-name: mixed-quickstart
    nodepool: nodepool-0
  name: test1
  namespace: default
spec:
  clusterName: mixed-quickstart
  replicas: 1
  selector:
    matchLabels:
      cluster.x-k8s.io/cluster-name: mixed-quickstart
      nodepool: nodepool-0
  template:
    metadata:
      labels:
        cluster.x-k8s.io/cluster-name: mixed-quickstart
        nodepool: nodepool-0
    spec:
      bootstrap:
        configRef:
          apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
          kind: KubeadmConfigTemplate
          name: test1-workers
        clusterName: mixed-quickstart
      infrastructureRef:
        apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
        kind: Metal3MachineTemplate
        name: test1-workers
        nodeDrainTimeout: 0s
        version: v1.30.0
---
apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
kind: Metal3MachineTemplate
metadata:
  name: test1-workers
  namespace: default
spec:
  template:
    spec:
      dataTemplate:
        name: test1-workers-template
      image:
        checksum: 7b2fbe69b2f2446d151b3e198b7fb020a4f17b9ab237c1b59b843e2783218b66
        checksumType: sha256
        format: qcow2
        url: http://192.168.222.1/UBUNTU_22.04_NODE_IMAGE_K8S_v1.30.0.qcow2
---
apiVersion: infrastructure.cluster.x-k8s.io/v1beta1
kind: Metal3DataTemplate
metadata:
  name: test1-workers-template
  namespace: default
spec:
  clusterName: mixed-quickstart
---
apiVersion: bootstrap.cluster.x-k8s.io/v1beta1
kind: KubeadmConfigTemplate
metadata:
  name: test1-workers
  namespace: default
spec:

```



```

template:
  spec:
    joinConfiguration:
      nodeRegistration:
        kubeletExtraArgs: {}
---
apiVersion: controlplane.cluster.x-k8s.io/v1beta1
kind: KubeadmControlPlane
metadata:
  name: mixed-quickstart-control-plane
  namespace: default
spec:
  kubeadmConfigSpec:
    clusterConfiguration:
      networking:
        dnsDomain: mixed-quickstart.default.local
        podSubnet: 192.168.0.0/16
        serviceSubnet: 10.128.0.0/16
      initConfiguration:
        nodeRegistration:
          criSocket: /var/run/containerd/containerd.sock
      joinConfiguration:
        nodeRegistration:
          criSocket: /var/run/containerd/containerd.sock
  machineTemplate:
    infrastructureRef:
      apiVersion: infrastructure.cluster.x-k8s.io/v1alpha1
      kind: KubevirtMachineTemplate
      name: mixed-quickstart-control-plane
      namespace: default
  replicas: 1
  version: v1.30.0
---
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha1
kind: KubevirtCluster
metadata:
  name: mixed-quickstart
  namespace: default
spec:
  controlPlaneEndpoint:
    host: 192.168.0.101
    port: 6443
  controlPlaneServiceTemplate:
    spec:
      type: LoadBalancer
---
apiVersion: infrastructure.cluster.x-k8s.io/v1alpha1
kind: KubevirtMachineTemplate
metadata:
  name: mixed-quickstart-control-plane
  namespace: default
spec:
  template:
    spec:
      virtualMachineBootstrapCheck:
        checkStrategy: ssh
      virtualMachineTemplate:
        metadata:
          namespace: default
        spec:
          runStrategy: Always
          template:
            spec:
              domain:
                cpu:
                  cores: 2
              devices:
                disks:

```

```
- disk:
  bus: virtio
  name: containervolume
  networkInterfaceMultiqueue: true
memory:
  guest: 4Gi
evictionStrategy: External
volumes:
- containerDisk:
  image: quay.io/capk/ubuntu-2204-container-disk:v1.30.1
  name: containervolume
```

Conclusion

This is not a perfect tutorial yet as we do not have real bare metal machines to finish all the setup, however, in the first phase, we aim to R&D the KubeVirt VM-related scenarios, which has been proven to be feasible in this doc. Once I get the bare metal machines, I will update this doc as soon as possible.