

前言: 问题描述

当前方案: 在多租户场景下, 交付以裸机 (目前主要指 x86) 为主要算力, KubeVirt VM为弹性算力的 k8s clusters。



问题描述: 由Management k8s cluster 管理和维护所有租户的 k8s clusters 的生命周期。

- 优点: 租户与租户之间是真实物理 k8s cluster 级别的隔离。
- 缺点: 1) 每个租户各自维护自己一套或者多套的 k8s clusters, 资源开销大; 2) 当租户数量达到一定数量后, 管理其 k8s clusters 会非常复杂。

解决方案 v1: 引入vCluster

前言: 解决方案 v2

愿景: 在vCluster的方案上继续提升隔离。

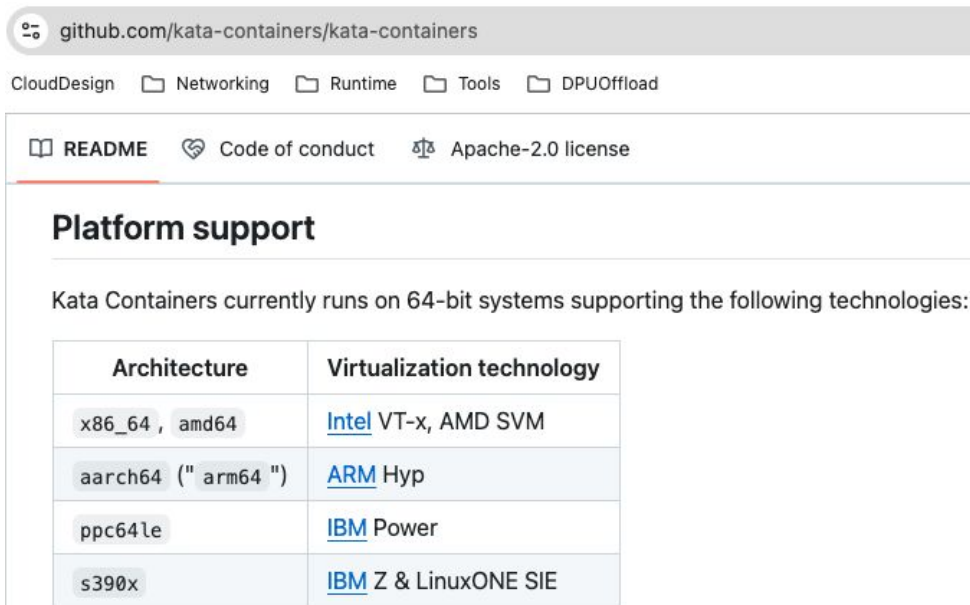


Agenda

- Kata Containers: Introduction & Features
- Kata Containers vs. Traditional Containers
- Demo: Install/Enable/Test Kata Containers in k8s
- Kata Containers' Key Components: Kata Containers Runtime & Kata Agent
- Kata Containers: Networking
- Kata Containers: Workflow
- **Integrate Kata Containers with vCluster**
- Reference

Kata Containers: Introduction

Kata Containers is an open source container runtime, building lightweight virtual machines that seamlessly plug into the containers ecosystem [1].



The screenshot shows the GitHub repository for Kata Containers. The breadcrumb navigation indicates the path: CloudDesign > Networking > Runtime > Tools > DPUOffload. The repository name is 'github.com/kata-containers/kata-containers'. Below the navigation bar, there are links for 'README', 'Code of conduct', and 'Apache-2.0 license'. The 'Platform support' section is highlighted with a red underline. It states that Kata Containers currently runs on 64-bit systems supporting the following technologies:

Architecture	Virtualization technology
x86_64 , amd64	Intel VT-x, AMD SVM
aarch64 (" arm64 ")	ARM Hyp
ppc64le	IBM Power
s390x	IBM Z & LinuxONE SIE

Note: Kata Containers is not supported on RISC-V yet.

Kata Containers: Features





 Security	Runs in a dedicated kernel, providing isolation of network, I/O and memory and can utilize hardware-enforced isolation with virtualization VT extensions.
 Compatibility	Supports industry standards including OCI container format, Kubernetes CRI interface, as well as legacy virtualization technologies.
 Performance	Delivers consistent performance as standard Linux containers; increased isolation without the performance tax of standard virtual machines.
 Simplicity	Eliminates the requirement for nesting containers inside full blown virtual machines; standard interfaces make it easy to plug in and get started.

Figure: Kata Containers' Features [2]

Kata Containers vs. Traditional Containers

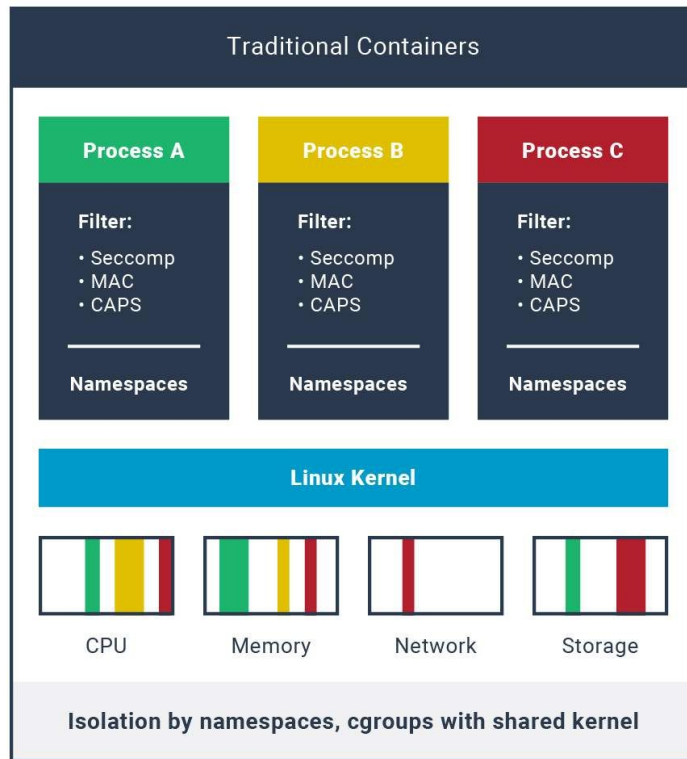
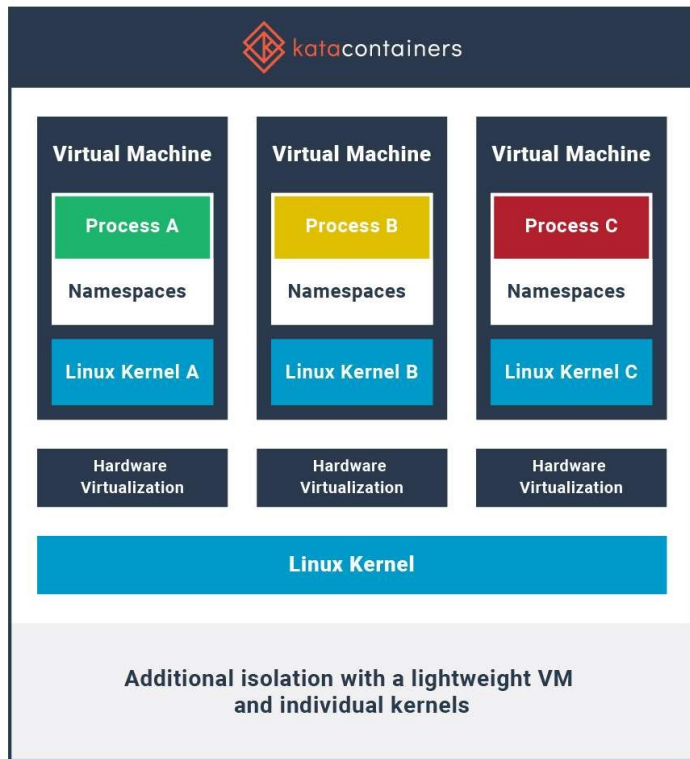


Figure: Kata Containers vs. Traditional Containers [2]

Demo: Install Kata Containers

```
root@master:~# kubectl get nodes -A
NAME        STATUS   ROLES    AGE   VERSION
master      Ready    control-plane   7d    v1.30.4
worker1     Ready    <none>        7d    v1.30.4

root@master:~# kubectl get pods -A -o wide
NAMESPACE   NAME                                                    READY   STATUS    RESTARTS   AGE   IP             NODE
default      php-apache-kata-qemu-689bf9f4b8-pjpnf                0/1     ImagePullBackOff    0      7m22s   10.244.235.154   worker1
kube-system  calico-kube-controllers-57cc879486-htv47             1/1     Running    13 (61m ago)   7d    10.244.219.69   master
kube-system  calico-node-97s6b                                     1/1     Running    4 (74m ago)   4d7h   192.168.122.200 worker1
kube-system  calico-node-fc7qq                                     1/1     Running    10 (61m ago)   4d7h   192.168.122.114 master
kube-system  coredns-7b5944fdcf-d7ljv                             1/1     Running    9 (61m ago)   7d    10.244.219.72   master
kube-system  coredns-7b5944fdcf-lpmh7                             1/1     Running    9 (61m ago)   7d    10.244.219.71   master
kube-system  etcd-master                                           1/1     Running    11 (61m ago)   7d    192.168.122.114 master
kube-system  kata-deploy-6zv65                                     1/1     Running    0          56m    10.244.219.73   master
kube-system  kata-deploy-ggppp                                     1/1     Running    0          56m    10.244.235.135   worker1
kube-system  kube-apiserver-master                               1/1     Running    13 (61m ago)   7d    192.168.122.114 master
kube-system  kube-controller-manager-master                     1/1     Running    14 (61m ago)   7d    192.168.122.114 master
kube-system  kube-proxy-gkmmn                                     1/1     Running    4 (74m ago)   7d    192.168.122.200 worker1
kube-system  kube-proxy-przqn                                     1/1     Running    10 (61m ago)   7d    192.168.122.114 master
kube-system  kube-scheduler-master                              1/1     Running    14 (61m ago)   7d    192.168.122.114 master
```

kubectl apply -f <https://raw.githubusercontent.com/kata-containers/kata-containers/main/tools/packaging/kata-deploy/kata-deploy/base/kata-deploy.yaml>

kubectl apply -f <https://github.com/kata-containers/kata-containers/blob/main/tools/packaging/kata-deploy/kata-rbac/base/kata-rbac.yaml>

Demo: Enable Kata Container RuntimeClass

```
root@master:~# kubectl get runtimeclasses
```

NAME	HANDLER	AGE
kata-clh	kata-clh	11m
kata-cloud-hypervisor	kata-cloud-hypervisor	11m
kata-dragonball	kata-dragonball	11m
kata-fc	kata-fc	11m
kata-qemu	kata-qemu	11m
kata-qemu-coco-dev	kata-qemu-coco-dev	11m
kata-qemu-nvidia-gpu	kata-qemu-nvidia-gpu	11m
kata-qemu-nvidia-gpu-snp	kata-qemu-nvidia-gpu-snp	11m
kata-qemu-nvidia-gpu-tdx	kata-qemu-nvidia-gpu-tdx	11m
kata-qemu-runtime-rs	kata-qemu-runtime-rs	11m
kata-qemu-se	kata-qemu-se	11m
kata-qemu-sev	kata-qemu-sev	11m
kata-qemu-snp	kata-qemu-snp	11m
kata-qemu-tdx	kata-qemu-tdx	11m
kata-remote	kata-remote	11m
kata-stratovirt	kata-stratovirt	11m

kubectl apply -f <https://raw.githubusercontent.com/kata-containers/kata-containers/main/tools/packaging/kata-deploy/runtimeclasses/kata-runtimeClasses.yaml>

Demo: Start a Pod with kata-qemu RuntimeClass

```
root@master:/home/test# kubectl get pods -o wide
NAME                                READY   STATUS    RESTARTS   AGE   IP              NODE       NOMINATED NODE   READINESS GATES
php-apache-kata-qemu-77b7cdcc9d-cjq8b 1/1     Running   0          2d15h  10.244.235.151  worker1    <none>           <none>

root@master:/home/test#
root@master:/home/test# kubectl describe pod php-apache-kata-qemu-77b7cdcc9d-cjq8b | grep "Runtime Class Name"
Runtime Class Name: kata-qemu
root@master:/home/test# kubectl describe pod php-apache-kata-qemu-77b7cdcc9d-cjq8b | grep "cni.projectcalico.org/containerID"
Annotations:      cni.projectcalico.org/containerID: b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e
root@master:/home/test#
root@master:/home/test# curl http://10.244.235.151:80; echo
OK!
```

```
root@worker1:/home/test# ps -ef | grep qemu
root      44761   44751   0 Sep06 ?           00:00:04 /opt/kata/bin/qemu-system-x86_64 -name sandbox-b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e -uuid 67c7bf6a-995c-4f11-bf3a-f22aa4b84bb2 -machine q35,accel=kvm,nvdimms=on -cpu host,pmu=off -qmp unix:fd=3,server=on,wait=off -m 2048M,slots=10,maxmem=8960M -device pci-bridge,bus=pcie.0,id=pci-bridge-0,chassis_nr=1,shpc=off,addr=2,io-reserve=4k,mem-reserve=1m,pref64-reserve=1m -device virtio-serial-pci,disable-modern=true,id=serial0 -device virtconsole,chardev=charconsole0,id=console0 -chardev socket,id=charconsole0,path=/run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/console.sock,server=on,wait=off -device nvdimms,id=nv0,memdev=mem0,unarmed=on -object memory-backend-file,id=mem0,mem-path=/opt/kata/share/kata-containers/kata-ubuntu-latest.image,size=268435456,readonly=on -device virtio-scsi-pci,id=scsi0,disable-modern=true -object rng-random,id=rng0,filename=/dev/urandom -device virtio-rng-pci,rng=rng0 -device vhost-vsock-pci,disable-modern=true,vhostfd=4,id=vsock-1227470637,guest-cid=1227470637 -chardev socket,id=char-bfd9b3be85dbc867,path=/run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/vhost-fs.sock -device vhost-user-fs-pci,chardev=char-bfd9b3be85dbc867,tag=kataShared,queue-size=1024 -netdev tap,id=network-0,vhost=on,vhostfds=5,fds=6 -device driver=virtio-net-pci,netdev=network-0,mac=96:66:40:48:86:7f,disable-modern=true,mq=on,vectors=4 -rtc base=utc,driftfix=slew,clock=host -global kvm-pit.lost_tick_policy=discard -vga none -no-user-config -nodefaults -nographic --no-reboot -object memory-backend-file,id=dimms1,size=2048M,mem-path=/dev/shm,share=on -numa node,memdev=dimms1 -kernel /opt/kata/share/kata-containers/vmlinuz-6.1.62-134 -append tsc=reliable no_timer_check rcupdate.rcu_expedited=1 i8042.direct=1 i8042.dumbkbd=1 i8042.nopnp=1 i8042.noaux=1 noreplace-smp reboot=k cryptomgr.no_tests net.ifnames=0 pci=lastbus=0 root=/dev/pmem0p1 rootflags=dax,data=ordered,errors=remount-ro ro rootfstype=ext4 console=hvc0 console=hvc1 quiet systemd.show_status=false panic=1 nr_cpus=8 selinux=0 systemd.unit=kata-containers.target systemd.mask=systemd-networkd.service systemd.mask=systemd-networkd.socket scsi_mod.scan=none -pidfile /run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/pid -smp 1,cores=1,threads=1,sockets=8,maxcpus=8
```

kubectl apply -f <https://raw.githubusercontent.com/kata-containers/kata-containers/main/tools/packaging/kata-deploy/examples/test-deploy-kata-qemu.yaml>

Kata Containers Runtime

The Kata Containers runtime is compatible with the OCI runtime specification and therefore works seamlessly with the Kubernetes Container Runtime Interface (CRI) through the CRI-O and containerd implementations.

Kata Containers provides a "shimv2" compatible runtime.

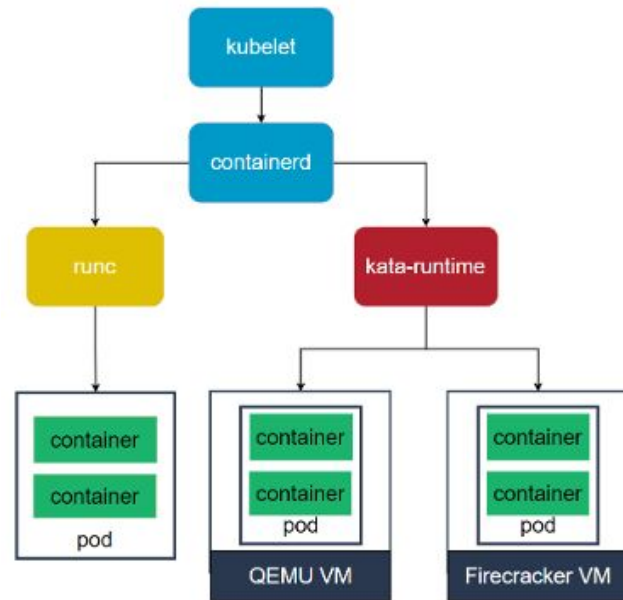



Figure: Integrate Kata Containers with k8s [3]

Kata Agent

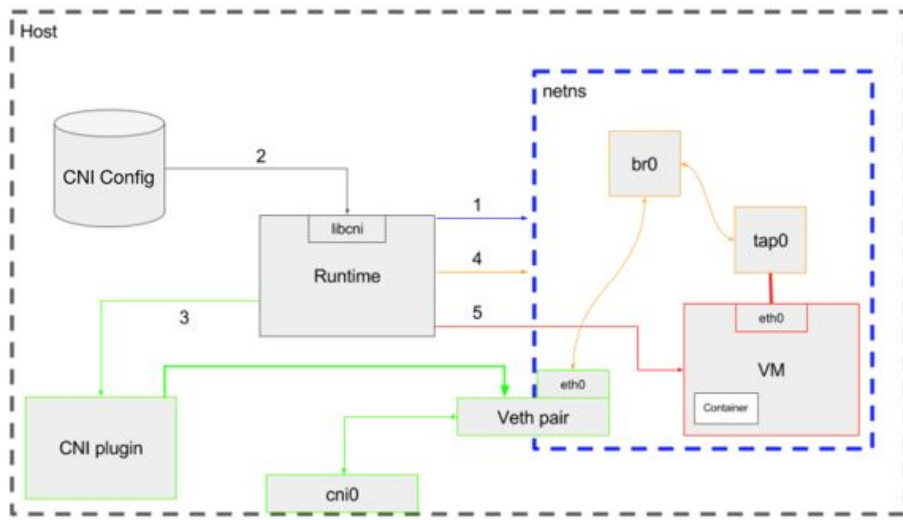


Kata containers manage the agent through a VSOCK exposed by the microVM to the host.
This VSOCK is used to communicate with the agent.

The agent is responsible for communicating between the microVM and the container workload.

```
grep qemu | sed 's/-device/\n-device/g'
?      00:00:09 /opt/kata/bin/qemu-system-x86_64 -name sandbox-b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e -u
4b84bb2 -machine q35,accel=kvm,nvdimms=on -cpu host,pmu=off -qmp unix:fd=3,server=on,wait=off -m 2048M,slots=10,maxmem=8960M
pci-bridge-0,chassis_nr=1,shpc=off,addr=2,io-reserve=4k,mem-reserve=1m,pref64-reserve=1m
-modern=true,id=serial0
console0,id=console0 -chardev socket,id=charconsole0,path=/run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/
0,unarmed=on -object memory-backend-file,id=mem0,mem-path=/opt/kata/share/kata-containers/kata-ubuntu-latest.image,size=268435456,read
disable-modern=true -object rng-random,id=rng0,filename=/dev/urandom
-device virtio-rng-pci,rng=rng0
-device vhost-vsock-pci,disable-modern=true,vhostfd=4,id=vsock-1227470637,guest-cid=1227470637 -chardev socket,id=char-bfd9b3be85dbc867,path=/run/vc/vm/b7666669041e29
734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/vhost-fs.sock
-device vhost-user-fs-pci,chardev=char-bfd9b3be85dbc867,tag=kataShared,queue-size=1024 -netdev tap,id=network-0,vhost=on,vhostfds=5,fds=6
-device driver=virtio-net-pci,netdev=network-0,mac=96:66:40:48:86:7f,disable-modern=true,mq=on,vectors=4 -rtc base=utc,driftfix=slew,clock=host -global kvm-pit.lost_t
ick_policy=discard -vga none -no-user-config -nodefaults -nographic --no-reboot -object memory-backend-file,id=dimm1,size=2048M,mem-path=/dev/shm,share=on -numa node,
memdev=dimm1 -kernel /opt/kata/share/kata-containers/vmlinuz-6.1.62-134 --append tsc=reliable no_timer_check rcupdate.rcu_expedited=1 i8042.direct=1 i8042.dumbkbd=1 i8
042.nopnp=1 i8042.noaux=1 noreplace-smp reboot=k cryptomgr.notests net.ifnames=0 pci=lastbus=0 root=/dev/pmem0p1 rootflags=dax,data=ordered,errors=remount-ro ro rootf
stype=ext4 console=hvc0 console=hvc1 quiet systemd.show_status=false panic=1 nr_cpus=8 selinux=0 systemd.unit=kata-containers.target systemd.mask=systemd-networkd.ser
vice systemd.mask=systemd-networkd.socket scsi_mod.scan=none -pidfile /run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/pid -smp 1,cores=1,t
hreads=1,sockets=8,maxcpus=8
```


Kata Containers: Networking

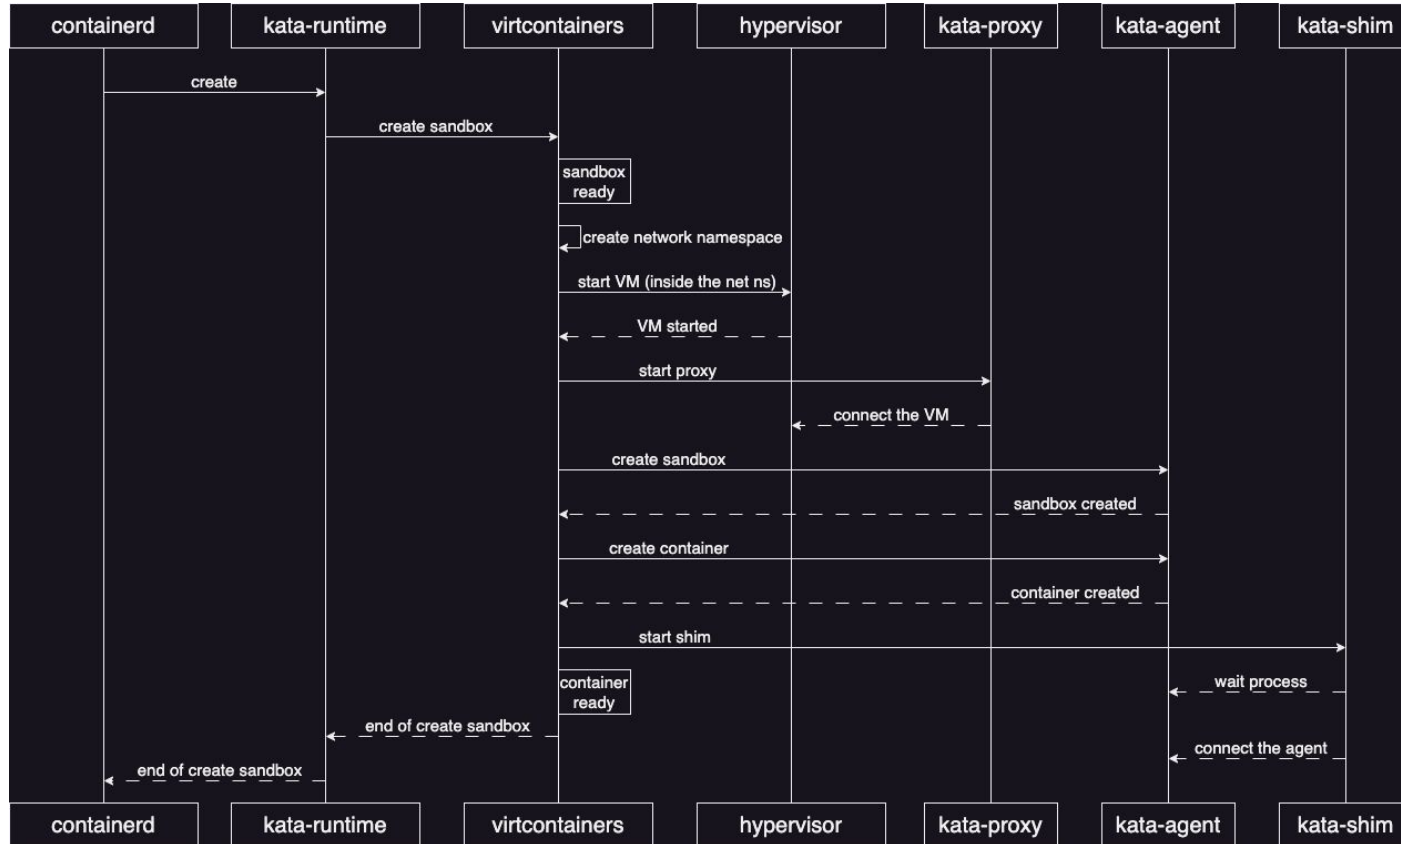


1. The runtime creates the blue-bordered network namespace, which contains all devices associated with the VM.
2. The runtime reads the required configuration from the CNI configuration files for the containers.
3. The runtime will communicate with the configured plug-in to start network service for the container. A veth pair is set up between cni0 and the container's network namespace.

```
4. A bridge inside the namespace is created, a tap device is
sandbox: b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e -u
f -qmp unix:fd=3,server=on,wait=off -m 2048M,slots=10,maxmem=8960M
mem-reserve=1m,pref64-reserve=1m
/run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/
namespace.
/opt/kata/share/kata-containers/kata-ubuntu-latest.image,size=268435456,read
dev/urandom
0637 -chardev socket,id=char-bfd9b3be85dbc867,path=/run/vc/vm/b7666669041e29
```

```
734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/vhost-fs.sock
Figure: Container Network Initiative (CNI) implementation for virtual-machine-based containers.
-device vhost-user-fs-pci,chardev=char-bfd9b3be85dbc867,tag=kataShared,queue-size=1024 -netdev tap,id=network-0,vhost=on,vhostfds=5,fds=6
-device driver=virtio-net-pci,netdev=network-0,mac=96:66:40:48:86:7f,disable-modern=true,mq=on,vectors=4 -rtc base=utc,driftfix=slew,clock=host -global kvm-pit.lost_t
ick_policy=discard -vga none -no-user-config -nodefaults -nographic --no-reboot -object memory-backend-file,id=dimm1,size=2048M,mem-path=/dev/shm,share=on -numa node,
memdev=dimm1 -kernel /opt/kata/share/kata-containers/vmlinux-6.1.62-134 -append tsc=reliable no_timer_check rcupdate.rcu_expedited=1 i8042.direct=1 i8042.dumbkbd=1 i8
042.nopnp=1 i8042.noaux=1 noreplace-smp reboot=k cryptomgr.notests net.ifnames=0 pci=lastbus=0 root=/dev/pmem0p1 rootflags=dax,data=ordered,errors=remount-ro ro rootf
stype=ext4 console=hvc0 console=hvc1 quiet systemd.show_status=false panic=1 nr_cpus=8 selinux=0 systemd.unit=kata-containers.target systemd.mask=systemd-networkd.ser
vice systemd.mask=systemd-networkd.socket scsi_mod.scan=none -pidfile /run/vc/vm/b7666669041e29734acac4f86b54870405b7f829b83e262d7b84bd719aa5d08e/pid -smp 1,cores=1,t
hreads=1,sockets=8,maxcpus=8
```

Kata Containers: Workflow



Demo: Kata Containers + vCluster

<https://github.com/loft-sh/vcluster/issues/2125>

KubeCon + CloudNative North America 2022 has ended

Running Isolated VirtualClusters With Kata & Cluster API - Chris Hein & Eric Ernst, Apple, Inc

📅 Wednesday October 26, 2022 3:25pm - 4:00pm EDT

📍 140 ABC



Running Isolated VirtualClusters With Kata & Cluster API - Chris Hein & Eric Ernst, Apple, Inc



BUILDING FOR THE ROAD AHEAD

DETROIT 2022

Running Isolated VirtualClusters With Kata & Cluster API

Chris Hein & Eric Ernst, Apple, Inc

前往平台观看:



References

- [1] Li, Guoqing & Takahashi, Keichi & Ichikawa, Kohei & Iida, Hajimu & Nakasan, Chawanat & Leelaprute, Pattara & Thiengburanathum, Pree & Phannachitta, Passakorn. (2023). The Convergence of Container and Traditional Virtualization: Strengths and Limitations. SN Computer Science. 4. 10.1007/s42979-023-01827-9.
- [2] <https://katacontainers.io/>
- [3] <https://github.com/kata-containers/documentation/blob/master/design/architecture.md>